

[University Name]
[Department Name]

*University Logo
Placeholder*

Autonomous Multi-Sensor Fire Detection Node Using Sensor Fusion and TinyML

Author Name

Supervisor: Prof. Dr. [Supervisor Name]

A Thesis Submitted in Partial Fulfillment of the Requirements for the Degree
of
Bachelor of Engineering

February 2026

Acknowledgments

First and foremost, I would like to express my deepest gratitude to my supervisor, [Supervisor's Name], for their invaluable guidance, patience, and constant encouragement throughout the development of this research. Their expertise in embedded systems and artificial intelligence, along with their insightful critiques, were instrumental in shaping the trajectory of this thesis.

I am also indebted to the technical staff at the Faculty of Engineering for providing the necessary laboratory resources and hardware support required for the experimental phase of this project. Special thanks are due to my peers and colleagues for the stimulating discussions and for providing a supportive environment that fostered critical thinking and technical rigor.

Furthermore, I wish to thank my family and friends for their unwavering support and belief in my abilities. Their encouragement was a source of strength during the most challenging phases of this academic journey.

Finally, I would like to acknowledge the developers and contributors of the open-source tools and platforms—specifically Arduino and Edge Impulse—whose technologies formed the foundation of the autonomous sensing node developed in this work.

Author Name
February 2026

Abstract

Traditional fire detection systems often struggle with false alarms triggered by non-combustion events such as cooking steam or aerosols. This thesis presents the development of an intelligent, autonomous multi-sensor fire detection node that leverages sensor fusion and TinyML to distinguish between genuine fires, ambient conditions, and common false alarm triggers. The system is built on the Arduino UNO R4 WiFi platform, utilizing its asymmetric multi-processing capabilities to decouple high-frequency sensor acquisition and local machine learning inference from low-priority telemetry tasks. A heterogeneous sensor suite—integrating smoke, volatile organic compounds (VOC), carbon monoxide (CO), infrared flame detection, and environmental temperature and humidity—provides a multi-modal “physical fingerprint” of the environment. A three-class classification model was developed and optimized using the Edge Impulse platform, employing an 8-bit quantized neural network for efficient edge deployment. Experimental results demonstrate that the system achieves 100% classification accuracy on a comprehensive validation dataset, with an on-device inference latency of less than 100 ms. A key finding of this research is the critical role of CO as a “truth sensor,” providing a reliable differentiator that remains stable during false alarm scenarios. The proposed architecture offers a robust, low-power, and high-fidelity solution for next-generation autonomous fire safety systems in smart building environments.

Table of Contents

1	Introduction	1
1.1	Background	1
1.2	Problem Statement	2
1.3	Research Questions	3
1.4	Research Objectives	3
1.5	Thesis Organization	4
2	Literature Review	4
2.1	Literature Selection Strategy	4
2.2	Evolution of Fire Detection and MEMS Technology	5
2.3	Sensor Fusion and TinyML at the Edge	6
2.4	False Alarm Mitigation and Research Synthesis	6
3	Theoretical Background	7
3.1	Physics and Operating Principles of Fire-Relevant Sensors	7
3.2	Characteristics of MEMS Smoke, VOC, CO, IR Flame, and Temperature/Humidity Sensors	8
3.3	Principles of Sensor Fusion for Discrimination	10
3.4	Feature Engineering for Fire Detection	11
3.4.1	Statistical Analysis (Mean, RMS, Kurtosis) for Gas/Thermal Trends	11
3.4.2	Frequency-Domain Features (FFT/Spectral Power) for Flame Flicker	12
3.4.3	Cross-Sensor Correlations and Dimensionality Reduction	12
3.5	Neural Network Fundamentals for Classification	13
3.5.1	The Artificial Neuron and Layered Architecture	14
3.5.2	Activation Functions: ReLU and Softmax	14
3.5.3	Training via Backpropagation and Gradient Descent	14
3.5.4	Non-Linear Multi-Class Discrimination	14
3.6	Model Quantization and Optimization for Resource-Constrained Devices	15
3.6.1	Principles of Neural Network Quantization	15
3.6.2	Performance Optimization and Latency Reduction	15

3.6.3	Pruning and Architectural Compression	16
3.7	Overview of the Edge Impulse Platform Pipeline	16
3.7.1	Data Ingestion and Labeling	16
3.7.2	Impulse Design and Digital Signal Processing (DSP)	16
3.7.3	Model Training and Validation	16
3.7.4	Deployment and Optimization	17
4	Sensor Selection and Characterization	17
4.1	Requirements for Multi-Modal Sensing	17
4.1.1	Functional and Physical Requirements	17
4.1.2	Operational and Integration Requirements	17
4.2	Detailed Examination of Selected Sensors	18
4.3	Justification of Sensor Choice	19
4.3.1	Sensitivity for Early Detection	19
4.3.2	Selectivity and the “Truth Sensor” Strategy	20
4.3.3	Redundancy and Contextual Validation	20
4.4	Sensor Calibration Protocols and Baseline Behavior	20
4.4.1	Hardware Pre-heating and Stabilization	20
4.4.2	Software-Level Baseline Detection	21
4.4.3	Flame Sensor Sensitivity Adjustment	21
4.4.4	Environmental Compensation (AHT20)	21
5	Hardware Platform and System Integration	21
5.1	The Autonomous Edge-Node Architecture	21
5.1.1	Justification for the Arduino UNO R4 WiFi Platform	22
5.1.2	Asymmetric Multi-Processing (AMP) Strategy	22
5.2	Power Management and Thermal Considerations	23
5.2.1	Power Demands of MEMS Gas Sensors	23
5.2.2	Thermal Management and Self-Heating	23
5.3	Physical Design and Enclosure	24
5.3.1	Airflow Considerations for Gas Sensing	24

5.3.2	Thermal and Modal Isolation	24
5.4	Connectivity and Location Awareness	25
5.4.1	MQTT Telemetry and Network Architecture	25
5.4.2	Static Location Tagging and Zone Identification	25
5.5	System Schematic and Wiring Diagrams	25
5.5.1	Digital Interconnectivity: The I2C Bus	26
5.5.2	Analog Interface and Pin Mapping	26
5.5.3	Power Distribution and Grounding	26
6	Data Collection Methodology	26
6.1	Safety Protocols During Data Collection	26
6.1.1	Hazard Identification and Risk Assessment	27
6.1.2	Fire Containment Measures	27
6.1.3	Ventilation Requirements	27
6.1.4	Personal Protective Equipment (PPE)	28
6.2	Rationale for Three-Class Classification	28
6.2.1	The Limitations of Binary Classification in Practical Fire Detection	28
6.2.2	Chemical and Thermal Overlap Between Nuisance and Fire Events	29
6.2.3	Three-Class Classification as a Design Solution	29
6.3	Sampling Parameters	29
6.3.1	Sampling Frequency	30
6.3.2	Window Size and Overlap Strategy	30
6.4	Class-Specific Data Collection Scenarios	31
6.5	Environmental Variance	38
6.5.1	Spatial and Contextual Variance	38
6.5.2	Rejection of Location-Specific Bias	39
6.6	Dataset Organization, Labeling, and Storage	39
7	Implementation	40
7.1	Edge Impulse Project Setup	40
7.1.1	Platform Rationale and Project Initialization	40

7.1.2	Data Ingestion and Labeling Configuration	41
7.1.3	Target Hardware Configuration and Resource Profiling	41
7.2	Impulse Design and DSP Blocks	42
7.2.1	Time-Series Windowing and Preprocessing	42
7.2.2	Spectral Analysis DSP Block	43
7.3	Model Architecture and Training	44
7.3.1	Neural Network Topology Design	44
7.3.2	Hyperparameter Tuning Results	45
7.4	Arduino Firmware Development	46
7.5	Heuristic Post-Processing and Hybrid Triggering Logic	49
7.6	Logic for Alarm Triggering	50
7.7	Quantization Strategy	51
7.7.1	Float32 vs. Int8: Precision, Memory, and Latency Trade-offs	51
7.7.2	Post-Training Quantization via Edge Impulse	52
8	Experimental Results and Evaluation	52
8.1	Test Environment Description	52
8.2	Dataset Summary	54
8.3	Training and Validation Metrics	56
8.4	Ablation Study	58
8.5	Environment-Specific Performance	61
8.6	On-Device Performance Metrics	62
8.7	Comparison with Baseline Approaches	64
8.7.1	Performance of Single-Modality Baselines	64
8.7.2	Advantages of Multi-Sensor Fusion	64
9	Discussion	65
9.1	Interpretation of Class Separability	65
9.1.1	The “CO Truth Sensor” and Combustion Verification	65
9.1.2	The “Heat Paradox” and Thermal Anomaly Analysis	65
9.1.3	Statistical Basis for Multi-Class Separability	66

9.2	Strengths of the “Autonomous Node” Approach	66
9.2.1	Localized Intelligence and Real-Time Inference	66
9.2.2	Resilience to Network Instability	66
9.2.3	High-Fidelity Discrimination via Multi-Sensor Fusion	67
9.2.4	Self-Awareness and Diagnostic Integrity	67
9.3	Analysis of Error Cases	67
9.3.1	The “Lab-to-Real-World” Gap	67
9.3.2	Sensor-Level Ambiguity and Fusion Limitations	67
9.3.3	Impact of Sensor Drift and Aging	68
9.4	Deployment Scenarios and Scalability	68
9.4.1	Cost Analysis per Node	68
9.4.2	Mesh Network Topology for Multi-Node Systems	69
9.4.3	Certification Challenges (UL/CE)	70
9.5	Ethical and Safety Aspects of AI in Life-Critical Systems	70
9.5.1	Transparency and the “Black Box” Problem	70
9.5.2	Bias and Data Integrity	71
9.5.3	Fail-Safe Design and Reliability	71
9.5.4	Accountability and Regulatory Compliance	71
10	Conclusion and Future Work	71
10.1	Limitations of This Study	71
10.1.1	Dataset Constraints	71
10.1.2	Lab vs. Real World	72
10.1.3	Sensor Drift	72
10.2	Summary of Key Findings	72
10.3	Contributions to the Field	73
10.3.1	Multi-Modal Sensor Fusion for High-Fidelity Discrimination	73
10.3.2	Optimized TinyML Deployment on Resource-Constrained Hardware	73
10.4	Future Work	74
10.4.1	Remote Model Management and Over-the-Air (OTA) Updates	74

10.4.2 Stress Testing and Noise-Injection for Reliability	74
10.4.3 Cooperative Sensing and Multi-Node Consensus	74
Bibliography	75

List of Figures

Figure 1	Sensor Correlation Matrix. A heatmap showing the Pearson correlation coefficients between the six sensors. Low correlation between certain sensors (e.g., CO and Humidity) indicates they provide unique, non-redundant information for the fusion model.	13
Figure 2	Indoor Fire Data Collection. Left: capturing infrared signatures from a flame source; Right: capturing combined flame and smoke signatures.	33
Figure 3	Smoldering Fire Simulation. The sensor node capturing particulate and VOC signatures from a smoldering source without visible flame.	34
Figure 4	Cooking Fume False Alarm Scenario. The sensor node capturing signatures from a frying pan, documenting the high VOC/Smoke but low CO profile.	36
Figure 5	Steam and Humidity False Alarm Capture. The sensor node capturing high humidity spikes from boiling water in both direct and ambient kitchen settings. .	37
Figure 6	Outdoor Night Data Collection. The sensor node capturing fire signatures in an outdoor environment to account for ambient noise and variance.	38
Figure 7	Edge Impulse Data Acquisition Interface. A screenshot of the platform’s data ingestion page, showing the raw time-series data from the multi-sensor array after being uploaded from the Arduino.	41
Figure 8	Impulse Design Architecture. The Edge Impulse “Create Impulse” screen, showing the configuration of the time-series data block, the Spectral Analysis DSP block, and the Neural Network Classifier.	42
Figure 9	Spectral Analysis DSP Block Configuration. Left: common errors during DSP block setup; Right: final optimized spectral analysis settings for multi-sensor fusion.	43
Figure 10	3D Feature Explorer Visualization. A plot showing how the three classes (fire, no_fire, false_alarm) cluster in the high-dimensional feature space, providing a visual proof of class separability.	44
Figure 11	Classifier Training Configuration. The training interface in Edge Impulse, detailing the neural network topology, learning rate, and training cycles (epochs) used for the fire detection model.	45
Figure 12	Live Model Testing and Validation. A screenshot of the model testing results within the Edge Impulse platform, showing the classification accuracy on the hold-out test dataset.	46
Figure 13	Deployment Target Selection. The deployment page showing the various export options, specifically the selection of the Arduino library for integration into the custom firmware.	48

Figure 14	Model Compilation and Build Success. The notification screen confirming the successful generation of the optimized C++ library for deployment on the Arduino UNO R4.	52
Figure 15	Dataset Class Distribution. A bar chart representing the number of samples for each of the three target classes: fire, no_fire, and false_alarm. The distribution confirms a well-balanced dataset, preventing model bias.	55
Figure 16	Model Confusion Matrix. A matrix visualizing the predicted vs. actual class labels. The 100% accuracy in this controlled environment demonstrates the high mathematical separability of the fire and false alarm classes under ideal conditions.	57
Figure 17	Sensor Signature Analysis by Class. Boxplots showing the median, quartiles, and outliers for each sensor across the three classes. This highlights the distinct chemical and thermal signatures of each scenario, such as higher VOC/Smoke in false alarms versus higher CO in real fires.	59
Figure 18	Feature Importance Ranking. A bar chart showing the relative importance of each sensor in the classification process. Smoke and VOC are identified as the most significant contributors (50% combined), while Humidity provides the least unique information.	59

List of Tables

Table 1	Comparative Analysis of Fire Detection Research Studies	7
Table 2	Summary of Selected Multi-Sensor Suite Specifications	18
Table 3	Functional Separation in the Arduino UNO R4 WiFi AMP Architecture	22
Table 4	Catalog of Data Collection Scenarios and Physical Signatures	32
Table 5	Estimated Bill of Materials (BOM) for the Autonomous Sensing Node	69

List of Abbreviations

AMP	Asymmetric Multi-Processing
CO	Carbon Monoxide
DSP	Digital Signal Processing
INT8	8-bit Integer
IR	Infrared
MEMS	Micro-Electro-Mechanical Systems
MQTT	Message Queuing Telemetry Transport
RAM	Random Access Memory
TinyML	Tiny Machine Learning
VOC	Volatile Organic Compounds

Glossary

ACH (Air Changes per Hour)	A measure of the air volume added to or removed from a space in one hour, divided by the volume of the space.
ADC (Analog-to-Digital Converter)	A system that converts an analog signal into a digital signal.
AMP (Asymmetric Multi-Processing)	A system architecture where multiple processors or cores operate independently, often with different roles or architectures.
ANN (Artificial Neural Network)	A computational model inspired by the structure and function of biological neural networks.
Adam (Adaptive Moment Estimation)	An optimization algorithm that can be used instead of the classical stochastic gradient descent procedure to update network weights iteratively based on training data.
Asymmetric Multi-Processing	A hardware architecture where different processor cores handle distinct tasks, such as separating high-speed inference from low-priority communication tasks.
Autonomous Sensing Node	A standalone edge device that performs data acquisition, processing, and decision-making locally without requiring a continuous cloud connection.
BOM (Bill of Materials)	An itemized list of all hardware components and their associated unit costs required to build one instance of the system.
Backpropagation	An algorithm used to calculate the gradient of a loss function with respect to the weights in a neural network.
Baseline	A standard or reference point used for comparison.
Baseline Resistance (\$R_0\$)	The electrical resistance of a gas sensor in clean air or a stable ambient environment.
Binary Classification	A type of classification task where there are only two possible classes or outcomes.
Blackbody Radiation	The electromagnetic radiation emitted by an idealized physical body that absorbs all incident electromagnetic radiation.
Burn-in Period	A period of time where a sensor is powered on and operated to stabilize its performance and remove initial contaminants.
Chemo-resistive Effect	The phenomenon where the electrical resistance of a material changes in response to the chemical adsorption of gas molecules.
Chimney Effect	The movement of air into and out of buildings, chimneys, or other containers, resulting from air buoyancy.

Complementary Fusion	A fusion strategy where sensors provide independent information about different aspects of the environment.
Convection	The transfer of heat through the movement of a fluid (such as air).
Cross-sensitivity	The sensitivity of a sensor to substances other than the target analyte.
DSP (Digital Signal Processing)	The use of digital processing, such as by computers or specialized digital signal processors, to perform a wide variety of signal processing operations.
Data Forwarder	A tool that allows for real-time streaming of sensor data from a device to the Edge Impulse studio.
Depletion Layer	An insulating region within a conductive semiconductor material where the mobile charge carriers have been diffused away.
Differential Sensing	A method of measurement where the change in a value relative to a baseline is more important than the absolute value.
EEPROM (Electrically Erasable Programmable Read-Only Memory)	A type of non-volatile memory used in computers and other electronic devices to store small amounts of data.
EON Compiler	Edge Impulse's deep learning compiler that optimizes models for low-power microcontrollers.
ESP-NOW	A connectionless, peer-to-peer WiFi communication protocol developed by Espressif for low-latency data transfer between ESP32 devices.
ESP32-S3	The secondary microcontroller on the Arduino UNO R4 WiFi responsible for WiFi and Bluetooth connectivity.
Edge AI	The practice of processing artificial intelligence algorithms locally on a hardware device, rather than on a remote cloud server.
Edge Deployment	The strategy of placing sensors and processing power at the "edge" of a network, directly at the site of data generation.
Edge Node	A computing element that performs data processing at the "edge" of the network, near the source of the data, to reduce latency and bandwidth use.
Environmental Compensation	The process of adjusting sensor readings to account for the influence of environmental factors like temperature and humidity.
Environmental Variance	The range of different environmental conditions (temperature, humidity, air quality) encountered by a system.

F1 Score	A measure of a model's accuracy on a dataset, calculated as the harmonic mean of precision and recall.
FDAS (Fire Detection and Alarm System)	A complete system comprising detectors, control panels, and notification appliances, assessed as a whole under safety standards.
FFT (Fast Fourier Transform)	An efficient algorithm for computing the Discrete Fourier Transform of a finite-length signal, decomposing it into constituent frequency components.
False Alarm Rate (FAR)	The frequency at which a system incorrectly indicates the presence of a target condition.
Field of View (FoV)	The extent of the observable world that is seen at any given moment by a sensor or camera.
Flame flicker frequency	The characteristic oscillation frequency (approximately 1–15 Hz) of turbulent diffusion flames, caused by periodic vortex shedding at the flame base.
Generalizability	The ability of a machine learning model to perform accurately on new, unseen data that was not part of its training set.
Ground Loop	An unwanted current path in a circuit that can cause signal interference and noise.
Heterogeneous Sensors	A collection of sensors that use different physical principles or target different phenomena.
I2C Bus	A multi-master, multi-slave, packet switched, single-ended, serial communication bus.
INT8	An 8-bit signed integer format.
IR-Transparent	A material that allows infrared radiation to pass through it with minimal absorption or reflection.
Impulse	The complete data processing and learning pipeline within the Edge Impulse platform.
Inference Latency	The time required for a deployed machine learning model to process an input and generate a classification output on the edge device.
JSON (JavaScript Object Notation)	A lightweight data-interchange format that is easy for humans to read and write and easy for machines to parse and generate.
Kurtosis	A fourth-order statistical moment measuring the heaviness of a distribution's tails; high kurtosis indicates impulsive, spike-like signal behaviour.

LDA (Linear Discriminant Analysis)	A supervised dimensionality reduction method that finds projections maximising the ratio of between-class to within-class variance.
Line-of-sight	A type of propagation that can transmit and receive data only where transmit and receive stations are in view of each other without any sort of an obstacle between them.
MEMS (Micro-Electro-Mechanical Systems)	Technology of very small devices, often used for high-sensitivity gas and pressure sensors on a single chip.
MLOps (Machine Learning Operations)	A set of practices that aims to deploy and maintain machine learning models in production reliably and efficiently.
MOS (Metal-Oxide Semiconductor)	A type of gas sensor that measures the change in resistance of a sensing layer when exposed to target gases.
MQTT (Message Queuing Telemetry Transport)	A lightweight, publish-subscribe network protocol that transports messages between devices.
Micro-heater	A miniaturized heating element used in MEMS sensors to maintain the sensing material at its required operating temperature.
Multi-Layer Perceptron (MLP)	A class of feedforward artificial neural network consisting of at least three layers of nodes.
Multi-class Classification	A classification task with more than two classes.
Multi-modal Sensing	A sensing strategy that utilizes multiple different types of sensors to gather information about an environment.
Notified Body	A conformity assessment organisation designated by an EU member state to evaluate products against harmonised European standards.
Nuisance Alarm	A false alarm triggered by non-hazardous environmental factors (e.g., steam, dust, or cooking) that mimics fire signatures.
Nuisance Event	An environmental event that may trigger a sensor but does not represent a genuine hazard, such as cooking steam or cigarette smoke.
Nuisance Source	A non-hazardous environmental stimulus that can trigger a sensor, such as cooking steam or cigarette smoke.
Nyquist–Shannon Sampling Theorem	A fundamental bridge between continuous-time signals and discrete-time signals, stating that a signal can be perfectly reconstructed if it is sampled at a rate greater than twice its highest frequency component.

Overlap	The portion of a sampling window that is shared with the previous or subsequent window.
PCA (Principal Component Analysis)	An unsupervised linear transformation that projects data onto orthogonal axes of maximum variance, reducing dimensionality while retaining most information.
PLS-DA (Partial Least Squares Discriminant Analysis)	A statistical method used to find a linear regression model by projecting the predicted variables and the observable variables to a new space.
PPE (Personal Protective Equipment)	Specialized clothing or equipment worn by employees for protection against health and safety hazards.
PRISMA	Preferred Reporting Items for Systematic Reviews and Meta-Analyses; a standardized evidence-based minimum set of items for reporting in reviews.
Post-Training Quantization (PTQ)	A technique to convert a pre-trained floating-point model to a quantized model with minimal accuracy loss and without retraining.
Publish-Subscribe	A messaging pattern where senders (publishers) do not send messages directly to specific receivers (subscribers), but instead characterize published messages into classes without knowledge of which subscribers there may be.
Pull-up Resistor	A resistor used to ensure a known state for a signal line.
Quantization	A model optimization technique that converts high-precision weights (e.g., 32-bit float) to lower-precision formats (e.g., 8-bit integer) to reduce memory usage and increase inference speed.
Quantization-Aware Training (QAT)	A method where the model is trained with quantization effects in the loop to improve accuracy in low-precision deployments.
RA4M1	The primary 32-bit ARM Cortex-M4 microcontroller on the Arduino UNO R4 WiFi.
RMS (Root Mean Square)	The square root of the mean of squared sample values, providing a measure of signal energy or effective amplitude over a window.
ReLU (Rectified Linear Unit)	A linear function that will output the input directly if it is positive, otherwise, it will output zero.
Redox Reaction	A chemical reaction involving the transfer of electrons between two species, consisting of simultaneous reduction and oxidation.
Redundant Fusion	A fusion strategy where multiple sensors provide independent measurements of the same physical property.

SIMD (Single Instruction, Multiple Data)	A type of parallel processing that allows one instruction to be applied to multiple data points simultaneously.
STEL (Short-Term Exposure Limit)	The maximum concentration of a chemical to which workers may be exposed continuously for a short period of time without any danger to health.
Sampling Frequency	The number of samples per unit time taken from a continuous signal to make a discrete signal.
Selectivity	The ability of a sensor to respond specifically to a target analyte while rejecting interference from other substances.
Self-heating	The phenomenon where an electronic device or sensor increases in temperature due to its own internal power dissipation.
Sensitivity	The ability of a sensor to detect small changes in the target analyte or physical property.
Sensor Fusion	The process of combining data from multiple sensors to achieve higher accuracy and more reliable information than possible with single sensors.
Smoke Dilution	The reduction in smoke concentration due to mixing with clean air, often caused by strong ventilation or wind.
Smoldering	A slow, low-temperature, flameless form of combustion, sustained by the heat evolved when oxygen directly attacks the surface of a condensed-phase fuel.
Softmax	An activation function that turns a vector of numbers into a vector of probabilities that sum to one.
Specificity	The ability of a test to correctly identify those without the condition (true negative rate).
Spectral Leakage	The phenomenon where signal energy at one frequency "leaks" into adjacent frequency bins in the FFT spectrum.
Spectral power	The squared magnitude of a signal's Fourier coefficients within a specified frequency band, representing energy concentration at those frequencies.
Static Location Tagging	Assignment of a fixed Zone_ID string to each node's MQTT payload at firmware configuration time.
Sub-ppm	Concentrations less than one part per million.
T90	The time required for a sensor to reach 90% of its final stable reading after a step change in concentration.
TWA (Time-Weighted Average)	The average exposure to a contaminant over a specified period, usually an 8-hour workday.

Telemetry	The automatic measurement and wireless transmission of data from remote sources.
Three-Class Classification	A supervised learning approach that categorizes data into three distinct labels—in this context, identifying fire, normal background, or non-fire nuisance events.
TinyML	A field of machine learning focused on deploying optimized models on low-power, resource-constrained microcontrollers at the hardware edge.
UL (Underwriters Laboratories)	A US-based safety certification organisation.
Veto Logic	A decision-making strategy where a negative signal from a highly reliable sensor (truth sensor) can override positive signals from other sensors to prevent a false alarm.
Voltage Rail	A single voltage provided by a power supply unit, such as 3.3V or 5V.
Welch Window	A windowing function used in spectral analysis to reduce the influence of signal discontinuities at the edges of a data segment.
Window Size	The duration of a signal segment used for a single inference or analysis.
Zone_ID	A unique identifier assigned to a sensing node that represents its physical location in a building.

1 Introduction

1.1 Background

This section establishes the context for the research by examining the global impact of fire incidents and the critical shortcomings of current fire safety infrastructure. It begins by discussing the high rates of false alarms that lead to occupant desensitization and significant economic losses. The discussion then contrasts legacy systems with the capabilities of intelligent edge nodes, highlighting how multi-sensor fusion and on-device machine learning can address the limitations of single-parameter sensors.

1.1.0.1 Significance of Fire Detection and False Alarm Issues

Fire-related incidents continue to pose a severe threat to global life safety, infrastructure, and ecological stability. In 2024, the United States reported approximately 1.39 million fires, which resulted in 3,920 civilian deaths and an estimated \$19.1 billion in direct property damage. Beyond the residential domain, wildland fires have intensified due to shifting climate patterns; the 2024–2025 fire season recorded carbon emissions of approximately 2.2 Pg C, which is 9% above the historical average. These extreme wildfire events, particularly in the North American boreal forests and South American wetlands, underscore the urgent need for autonomous, rapid-response detection systems capable of operating in diverse environmental conditions.

However, the efficacy of modern fire safety infrastructure is severely undermined by the prevalence of false alarms, often referred to as nuisance alarms. In the United Kingdom, Home Office data for the 2020/21 period revealed that 98% of all incidents triggered by legacy systems were false alarms, with only 2% resulting from actual fires. Approximately 90% of these false positives are attributed to faulty apparatus or environmental interference rather than malicious intent. This “cry wolf” effect leads to dangerous occupant desensitization, where individuals may ignore or disable life-saving equipment due to frequent erroneous triggers.

The economic ramifications of these inaccuracies are profound. In the UK, false fire alarms cause annual losses exceeding £1 billion due to lost productivity and business disruption, with the average cost per instance estimated at £2,900. Similarly, regional reports from New South Wales, Australia, indicate that false alarms account for nearly 37% of total fire department responses, straining public resources and delaying deployment to genuine emergencies. Consequently, reducing the false alarm rate is a critical requirement for maintaining the integrity of fire services and public trust in safety technology.

1.1.0.2 Limitations of Legacy Systems vs. Intelligent Edge Nodes

Legacy fire detection systems primarily rely on single-parameter sensing, such as ionization or photoelectric smoke detection. While effective in specific scenarios, these sensors lack the granularity to distinguish between genuine combustion and common household or industrial non-fire aerosols. Photoelectric sensors are notoriously prone to interference from water vapor (steam), dust, and cooking fumes, while ionization sensors frequently trigger in response to high humidity or minor culinary activities. These single-modality systems operate on static thresholding, which fails to account for the complex, time-varying signatures of early-stage fires or the cross-correlation between temperature, gas concentrations, and light radiation.

The transition toward intelligent edge nodes—also termed *Autonomous Sensing Nodes*—offers a paradigm shift in detection reliability. By integrating a multi-sensor suite—incorporating Volatile Organic Compounds (VOC), Carbon Monoxide (CO), Infrared (IR) flame detection, and environmental variables (temperature and humidity)—a node can capture the multi-dimensional fingerprint of a fire event. Modern microcontrollers, such as the Arduino UNO R4 WiFi, feature an *Asymmetric Multi-Processing* architecture that combines a high-performance Renesas Cortex-M4 core with a dedicated ESP32-S3 co-processor for connectivity. This setup is ideal for deploying Tiny Machine Learning (TinyML) models directly at the sensing site. It allows for real-time sensor fusion where the system processes high-frequency data (e.g., 10 Hz) across temporal windows (typically 10 s) to identify patterns that static thresholds overlook.

Unlike traditional sensors, these intelligent nodes utilize local inference to classify environments into three distinct states: `fire`, `no_fire`, and `false_alarm`. This classification framing significantly reduces latency by eliminating the need for cloud-based processing while preserving battery life through efficient on-device computation. Research indicates that TinyML-driven multi-sensor fusion can achieve detection accuracies exceeding 98% while concurrently reducing false alarm rates to below 1%, effectively overcoming the physical limitations of legacy hardware.

1.2 Problem Statement

1.2.0.1 Background and Problem Statement

The prevalence of false alarms in traditional fire detection systems remains a critical challenge for emergency response efficiency and public safety trust. Legacy detectors, primarily relying on single-parameter sensing like ionization or photoelectric smoke detection, often fail to distinguish between genuine combustion events and common environmental “nuisance” triggers such as cooking steam, aerosols, or cleaning vapors. This lack of granularity leads to significant economic loss and “alarm fatigue,” where occupants may ignore or even disable life-safety systems. Consequently, there is an urgent need for intelligent, multi-sensor detection platforms capable of local, high-fidelity discrimination between fire and non-fire conditions.

Furthermore, while multi-modal sensing offers a potential solution, the requirement for high-frequency sampling and real-time processing historically necessitated centralized computing architectures. Such architectures introduced unacceptable latencies and relied on persistent network connectivity, which can be compromised during a fire. The problem, therefore, is not only to improve detection accuracy but to do so within the constraints of an autonomous, edge-deployed sensing node that can perform complex sensor fusion and machine learning inference locally on a low-power microcontroller.

1.2.0.2 Scope and Limitations

This research focuses on the development and evaluation of a stationary autonomous sensing node based on the Arduino UNO R4 WiFi platform. The scope is limited to the detection of incipient fire stages and the mitigation of common indoor false alarm triggers through TinyML-driven multi-sensor fusion.

The system’s “Autonomous” nature refers to its capacity for local data acquisition, feature extraction, and classification inference without reliance on cloud-based processing for primary

life-safety decisions. While the node includes MQTT telemetry for remote monitoring and reporting, its core detection logic is executed entirely on the Renesas RA4M1 core.

It is important to distinguish this “Autonomous Sensing Node” from a mobile robotic detector. The node is designed for stationary deployment in residential, commercial, or industrial settings. Its fixed position allows for the establishment of stable environmental baselines, which is a key component of its discrimination logic. The research does not cover the navigation, obstacle avoidance, or active suppression capabilities required for mobile fire-fighting robots. Furthermore, while the node can operate on battery power, long-term deployment strategies for energy harvesting or extreme low-power sleep modes are considered secondary to the primary objective of false alarm reduction through sensor fusion.

1.3 Research Questions

This research aims to evaluate the effectiveness and feasibility of an autonomous sensing node for fire detection. To guide the investigation, the following research questions have been formulated:

1. **Can a multi-sensor fusion approach distinguish between real fire, ambient conditions, and common false alarm triggers?** This question explores the effectiveness of combining five heterogeneous sensor modalities (smoke, VOC, CO, IR flame, and environmental context) to create a physically-grounded “fingerprint” of combustion. It investigates whether the cross-correlation between different physical and chemical phenomena provides sufficient information to achieve high separability between classes that are often confused by single-modality systems.
2. **Is TinyML deployment on a Cortex-M4 microcontroller feasible for real-time fire detection with < 100ms latency?** This question addresses the technical feasibility of performing complex machine learning inference at the edge. It examines whether an optimized and quantized neural network model can operate within the severe memory and processing constraints of the Renesas RA4M1 core while maintaining the high sampling frequency (10 Hz) necessary for rapid fire response.
3. **How does three-class classification (fire/no_fire/false_alarm) compare to binary classification in reducing false positives?** This question investigates the architectural decision to explicitly include a “false_alarm” class in the model training. It seeks to determine if providing the model with labeled examples of common nuisance triggers (e.g., cooking fumes, alcohol vapors) results in a more robust decision boundary and a lower rate of spurious alerts compared to traditional binary models.

1.4 Research Objectives

The central objective of this research is to design, implement, and evaluate an autonomous sensing node that leverages multi-sensor fusion and TinyML for intelligent fire detection. The specific objectives are:

1. **To develop a multi-modal hardware platform** based on the Arduino UNO R4 WiFi that integrates MEMS sensors for smoke, VOC, CO, and AHT20 for environmental monitoring, along with an analog IR flame sensor.
2. **To implement an Asymmetric Multi-Processing (AMP) firmware architecture** that decouples high-frequency sensor acquisition and local TinyML inference (RA4M1) from low-priority MQTT telemetry (ESP32-S3).

3. **To collect and label a comprehensive multi-class dataset** encompassing genuine fire scenarios, ambient environmental conditions, and diverse false alarm triggers.
4. **To train and optimize a quantized neural network** (TinyML) using the Edge Impulse platform, capable of three-class classification with minimal latency on the target hardware.
5. **To evaluate the system's performance** in terms of classification accuracy, false alarm rejection, and on-device resource utilization (latency, RAM, Flash).

Through the fulfillment of these objectives, this research contributes a validated framework for the next generation of autonomous, high-fidelity fire detection nodes suitable for edge deployment in smart building ecosystems.

1.5 Thesis Organization

This thesis is structured into ten chapters, following the logical progression of the research from theoretical foundations to implementation and evaluation:

- **Chapter 1: Introduction** provides the research context, problem statement, research questions, and objectives.
- **Chapter 2: Literature Review** examines existing work in fire detection, sensor fusion, and TinyML, identifying the gaps this research seeks to address.
- **Chapter 3: Theoretical Background** outlines the physics of the sensors used and the mathematical principles behind sensor fusion and neural networks.
- **Chapter 4: Sensor Selection and Characterization** details the technical rationale for choosing specific sensors and their performance characteristics.
- **Chapter 5: Hardware Platform and System Integration** describes the autonomous node architecture and the integration of components on the Arduino UNO R4 WiFi.
- **Chapter 6: Data Collection Methodology** explains the experimental setup, safety protocols, and the strategy for creating a labeled multi-class dataset.
- **Chapter 7: Implementation** details the Edge Impulse pipeline, neural network training, and the development of the Arduino firmware.
- **Chapter 8: Experimental Results and Evaluation** presents the classification performance, ablation studies, and on-device resource metrics.
- **Chapter 9: Discussion** interprets the findings, discusses the strengths and limitations of the approach, and addresses ethical and safety considerations.
- **Chapter 10: Conclusion and Future Work** summarizes the key contributions and proposes avenues for further research.

2 Literature Review

2.1 Literature Selection Strategy

The comprehensive literature review undertaken for this thesis systematically examined scholarly works published between 2015 and the present, ensuring an up-to-date understanding of advancements in intelligent fire detection systems. The primary focus of this selection strategy was threefold: multi-modal sensing, the application of TinyML for edge deployment, and innovative approaches to false alarm reduction. This targeted approach aimed to identify research that addresses the limitations of conventional fire detection systems, particularly their susceptibility to false positives.

The evolution of fire detection research has seen a significant shift from traditional threshold-based systems to more sophisticated, data-driven methodologies (Alajlan & Ibrahim, 2022). Early systems often relied on single-sensor thresholds, leading to high rates of false alarms caused by environmental interferences such as cooking fumes or steam (Fonollosa et al., 2018). To mitigate these issues, the current literature emphasizes the integration of multiple sensor modalities. Studies have explored combining temperature, smoke, flame, and carbon monoxide (Solórzano & others, 2021). The consistent identification of CO as a crucial differentiator between actual combustion events and false alarms in multi-sensor setups underlines the importance of multi-modal sensing strategies.

A critical aspect of the literature search involved identifying research pertaining to TinyML, particularly its application in resource-constrained edge devices like Arduino platforms. This area of study is vital for developing autonomous sensing nodes capable of real-time inference with minimal latency (Makkapati et al., 2022). The integration of TinyML enables complex machine learning models to run directly on microcontrollers.

Furthermore, the literature selection prioritized studies focused on reducing false alarms, a pervasive challenge in fire detection. Research has moved beyond merely differentiating fire from non-fire conditions to more nuanced classification, including the explicit recognition of false alarm scenarios. This reflects a growing understanding that effective fire detection systems must not only reliably identify true fires but also robustly distinguish them from common nuisance sources like burnt toast, steam, and aerosols, which often present with similar sensor signatures (Vorwerk et al., 2024).

2.2 Evolution of Fire Detection and MEMS Technology

The technological trajectory of fire detection has evolved from simple thermal triggers to sophisticated multi-modal sensing platforms. Early fire safety relied on ionization and photoelectric sensors, which, while effective at identifying smoke particles, lacked the intelligence to contextualize the environment. The primary limitation of these legacy systems is their reliance on static thresholding of a single physical parameter, making them highly susceptible to false alarms from non-fire aerosols, such as cooking steam or dust (Fonollosa et al., 2018).

The advent of Micro-Electro-Mechanical Systems (MEMS) has fundamentally changed this landscape. MEMS technology allows for the miniaturization of complex chemical and physical sensors, enabling the integration of multiple sensing modalities into a single, compact node. Modern metal-oxide (MOS) sensors for Smoke, Volatile Organic Compounds (VOCs), and Carbon Monoxide (CO) offer high sensitivity and rapid response times within a footprint of only a few millimeters (Pathan & others, 2024). These sensors operate by measuring changes in electrical resistance as gas molecules are adsorbed onto a heated sensing layer, a process that can be precisely controlled and monitored by microcontrollers.

Furthermore, the integration of Infrared (IR) flame detection adds an optical dimension to the sensing suite. Unlike gas sensors, which detect the chemical byproducts of combustion, IR sensors identify the radiant emission of an open flame, specifically the characteristic flicker in the 760 nm to 1100 nm band (Toreyin et al., 2012). The evolution toward “Autonomous Sensing Nodes” leverages these MEMS advancements to create distributed intelligence units that can capture a comprehensive “physical fingerprint” of a fire event, providing the multi-dimensional data necessary for robust discrimination between genuine fires and environmental noise (Meleti & Razi, 2024).

2.3 Sensor Fusion and TinyML at the Edge

The challenge of reliable fire detection is increasingly being addressed through the convergence of sensor fusion and Tiny Machine Learning (TinyML). Sensor fusion refers to the integration of data from multiple heterogeneous sensors to improve classification accuracy and robustness beyond what is possible with single-modality systems. In the context of fire detection, this involves identifying the cross-correlations between gas concentrations, IR radiation, and environmental variables (Pathan & others, 2024).

Traditional fusion techniques often relied on hand-crafted heuristic rules or centralized processing. However, the rise of TinyML allows for the deployment of complex neural networks directly on low-power microcontrollers, such as those based on the ARM Cortex-M architecture. This “Edge AI” approach eliminates the latency and privacy concerns associated with cloud-based processing, enabling autonomous nodes to make life-safety decisions locally and in real-time (Makkapati et al., 2022).

Research has demonstrated that quantized neural networks can achieve detection accuracies exceeding 98% while operating within the severe memory and computational constraints of embedded hardware (Alajlan & Ibrahim, 2022). A critical aspect of this workflow is model optimization, where techniques such as Post-Training Quantization (PTQ) are used to convert 32-bit floating-point weights into 8-bit integers, significantly reducing the model’s footprint with minimal impact on accuracy (Novac & others, 2021). By processing multi-sensor data at the edge, these intelligent nodes can identify the unique “fingerprint” of fire and reject localized false alarm triggers with high fidelity (Meleti & Razi, 2024).

2.4 False Alarm Mitigation and Research Synthesis

The mitigation of false alarms remains the primary driver for innovation in fire detection systems. Current literature identifies a significant gap in the ability of traditional detectors to distinguish between genuine combustion and common indoor nuisance events, such as cooking aerosols, water steam, and alcohol-based cleaning vapors. These events often produce physical signatures—elevated particulate counts and VOC spikes—that mimic the early stages of a fire, leading to high rates of spurious triggers (Fonollosa et al., 2018).

To address this challenge, recent research has moved toward three-class classification schemas (fire, no_fire, false_alarm) rather than binary detection. This approach involves training machine learning models on specific datasets that include labeled examples of nuisance triggers, allowing the classifier to learn the subtle differences between, for example, a cooking-induced smoke plume and actual wood combustion (Vorwerk et al., 2024).

Research Study	Sensors Used	Classification Type	Primary Innovation	Reported Accuracy
Fonollosa et al. (2018) (Fonollosa et al., 2018)	Chemical/ Gas	Binary	Early detection of chemical fires	92.4%
Liu et al. (2023) (Liu et al., 2023)	Multi-modal	3-Class	Identification of CO as “Truth Sensor”	98.1%
Meleti et al. (2024) (Meleti & Razi, 2024)	Thermal/ Optical	Binary	Obscured fire detection via IR	95.0%
This Work (2026)	6-Sensor Fusion	3-Class	TinyML False Alarm Recognition	100% (Lab)

Table 1: Comparative Analysis of Fire Detection Research Studies

Synthesis of the reviewed research underscores the importance of the Carbon Monoxide (CO) sensor as a “truth sensor” for combustion verification. While VOC and Smoke sensors are highly sensitive, they are also prone to cross-sensitivity with non-fire vapors; however, significant CO spikes are rarely present in common household nuisance scenarios, providing a reliable differentiator (Liu et al., 2023). The consensus in the literature points toward autonomous, multi-sensor nodes as the most promising solution for achieving high sensitivity and low false-alarm rates in complex indoor environments (Meleti & Razi, 2024).

3 Theoretical Background

3.1 Physics and Operating Principles of Fire-Relevant Sensors

The detection of incipient fire events requires the precise transduction of physical and chemical phenomena into measurable electrical signals. This process is governed by the operating principles of diverse sensor modalities, each targeting specific signatures of combustion, such as aerosol concentration, gas evolution, infrared radiation, and thermal fluctuations. Understanding these underlying physics is essential for developing robust sensor fusion algorithms that can distinguish between genuine fire threats and environmental noise.

3.1.0.1 Metal-Oxide Semiconductor (MOS) Gas Sensing

The primary mechanism for detecting smoke, volatile organic compounds (VOCs), and carbon monoxide (CO) in modern Micro-Electro-Mechanical Systems (MEMS) sensors is based on the chemo-resistive effect of metal-oxide semiconductors, most commonly tin dioxide (SnO₂). In an oxygen-rich environment, oxygen molecules from the atmosphere are adsorbed onto the surface of the heated metal-oxide grain, capturing electrons from the conduction

band and forming a depletion layer (Liu et al., 2023). This results in a high baseline electrical resistance.

When reducing gases—such as CO or hydrocarbons found in smoke—interact with the adsorbed oxygen, a redox reaction occurs, releasing electrons back into the semiconductor’s conduction band. This process reduces the thickness of the depletion layer and significantly decreases the sensor’s electrical resistance (Pathan & others, 2024). The magnitude of this resistance change is proportional to the concentration of the target gas, allowing for quantitative measurement. MEMS implementations optimize this process by integrating a micro-heater that maintains the sensing layer at an ideal operating temperature (typically 200°C to 400°C), ensuring rapid reaction kinetics and high sensitivity within a miniature footprint (Meleti & Razi, 2024).

3.1.0.2 Infrared (IR) Flame Detection

Flame detection relies on the principles of blackbody radiation and the specific spectral emissions of hot gases. During the combustion of organic materials, carbon dioxide (CO₂) molecules are excited and emit characteristic radiation in the infrared spectrum, particularly a strong peak at approximately 4.3 μm , known as the “CO₂ spike” (Deng & others, 2023). Additionally, the flickering nature of a flame—typically occurring at frequencies between 1 Hz and 20 Hz—provides a temporal signature that distinguishes it from static heat sources.

Infrared sensors, such as the Gravity Analog Flame Sensor, utilize a photodiode or phototransistor sensitive to a specific range of the IR spectrum (typically 760 nm to 1100 nm). When IR photons within this band strike the active area of the sensor, they generate a photocurrent proportional to the radiation intensity (Toreyin et al., 2012). By analyzing both the absolute intensity and the frequency components of this signal, the system can identify the presence of an active flame while rejecting interference from sunlight or artificial lighting.

3.1.0.3 Capacitive and Resistive Environmental Sensing

Environmental variables, specifically temperature and humidity, serve as critical context for gas sensor readings and as secondary indicators of fire. Humidity sensing in modern digital sensors, such as the AHT20, often employs a capacitive operating principle. A thin-film polymer dielectric is placed between two electrodes; as water molecules are adsorbed by the polymer, its dielectric constant changes, leading to a measurable change in capacitance (Liu et al., 2023).

Temperature sensing typically leverages the predictable resistance-temperature relationship of a thermistor or the voltage-temperature relationship of a p-n junction. In integrated MEMS devices, these thermal changes are converted into digital values via an on-chip Analog-to-Digital Converter (ADC), providing high-resolution data for environmental compensation and thermal anomaly detection (Meleti & Razi, 2024).

3.2 Characteristics of MEMS Smoke, VOC, CO, IR Flame, and Temperature/Humidity Sensors

The effectiveness of an autonomous fire detection node is fundamentally limited by the performance characteristics of its constituent sensors. In this research, a heterogeneous array of Micro-Electro-Mechanical Systems (MEMS) and analog sensors is employed to capture the multi-dimensional signature of fire and false alarm events. This section details the specific

characteristics, performance metrics, and inherent limitations of the sensors used in the autonomous node, including smoke, volatile organic compounds (VOC), carbon monoxide (CO), infrared (IR) flame, and environmental (temperature and humidity) sensors.

3.2.0.1 MEMS Gas Sensors (Smoke, VOC, and CO)

The gas sensing suite utilizes the DFRobot Fermion MEMS series, which represents a significant advancement over traditional electrochemical and bulky metal-oxide sensors. These MEMS devices are characterized by their extremely small footprint (typically 13 mm x 13 mm), low power consumption, and high sensitivity to target gases.

3.2.0.1.1 Smoke and VOC Detection (SEN0570 & SEN0566)

The MEMS smoke sensor (SEN0570) and VOC sensor (SEN0566) both utilize a metal-oxide semiconductor (MOS) sensing layer optimized for specific gas groups. The smoke sensor is particularly sensitive to the larger particulate matter and hydrocarbons typical of wood and paper combustion. A key characteristic of these sensors is their rapid response time, often achieving a T90 (90% of final value) within 30 seconds of exposure (DFRobot, 2024a).

However, these sensors exhibit significant cross-sensitivity; for instance, the smoke sensor is highly reactive to ethanol vapors, a property that is both a strength for detecting certain false alarms and a challenge for specificity (Pathan & others, 2024).

The VOC sensor targets a broader range of organic compounds, including formaldehyde, toluene, and benzene. Its sensing range typically spans from 0 to 1000 ppm, with a resolution capable of detecting sub-ppm changes in indoor air quality. Both sensors incorporate an internal micro-heater, requiring a stabilized pre-heating period (typically 48 hours for new sensors and 1-2 minutes after short power-off cycles) to achieve a stable baseline resistance (DFRobot, 2024b).

3.2.0.1.2 Carbon Monoxide Detection (SEN0564)

The MEMS CO sensor (SEN0564) is optimized for the detection of CO, a critical marker of incomplete combustion. Unlike the broad-spectrum VOC sensor, the CO sensor's sensing layer is tailored to be highly selective for CO molecules, with a typical detection range of 1 to 1000 ppm (DFRobot, 2024c). In the context of fire detection, this sensor serves as a "truth sensor," as elevated CO levels are rarely present in common non-fire scenarios such as cooking steam or aerosol usage, providing a vital differentiator for the sensor fusion model (Liu et al., 2023).

3.2.0.2 Infrared (IR) Flame Sensor (DFR0076)

The Gravity Analog Flame Sensor (DFR0076) is a phototransistor-based device designed to detect radiation in the 760 nm to 1100 nm wavelength range. This band corresponds to the infrared emissions of a typical flame. The sensor features a wide detection angle of approximately 60 degrees and a sensitivity that can be adjusted via an onboard potentiometer. Its response time is nearly instantaneous (< 1 ms), allowing for the capture of the high-frequency flicker characteristic of open flames (Toreyin et al., 2012). A significant limitation of this sensor modality is its line-of-sight requirement and susceptibility to intense IR interference from sunlight or incandescent lighting, necessitating its fusion with non-optical gas sensors to reduce false positives (Meleti & Razi, 2024).

3.2.0.3 Environmental Monitoring (AHT20)

The AHT20 sensor provides digital temperature and humidity data via the I2C protocol. It is characterized by its high accuracy ($\pm 0.3^\circ\text{C}$ for temperature and $\pm 2\%$ for relative humidity) and long-term stability.

- **Temperature:** The sensor operates across a -40°C to $+85^\circ\text{C}$ range. In fire detection, it monitors the “thermal anomaly” or “heat paradox,” where certain false alarms (like cooking) may actually show higher localized heat than incipient fires (Meleti & Razi, 2024).
- **Humidity:** The capacitive humidity sensor is sensitive to the 0-100% RH range. Humidity data is essential for compensating the cross-sensitivity of MOS gas sensors, as water vapor can occupy active sites on the sensing layer and influence resistance readings (Pathan & others, 2024).

3.3 Principles of Sensor Fusion for Discrimination

Sensor fusion is the process of integrating data from multiple heterogeneous sensors to produce information that is more accurate, reliable, and comprehensive than that provided by any single sensor in isolation. In the context of fire detection, sensor fusion addresses the inherent ambiguities of individual modalities—such as the susceptibility of photoelectric sensors to steam or the line-of-sight limitations of infrared sensors—by identifying cross-sensor correlations that characterize a genuine fire event.

3.3.0.1 Hierarchical Levels of Fusion

Sensor fusion can be categorized into three hierarchical levels based on the stage at which the data is combined: data-level (early fusion), feature-level, and decision-level (late fusion).

3.3.0.1.1 Data-Level and Feature-Level Fusion

Data-level fusion involves the direct integration of raw sensor signals. This approach preserves the highest degree of information but requires significant bandwidth and computational resources, as the fusion center must process high-dimensional raw data. Feature-level fusion, which is the primary approach used in this research, involves extracting relevant characteristics (features) from each sensor modality—such as mean gas concentration, thermal gradients, or the frequency components of a flame signal—before combining them into a single feature vector (Pathan & others, 2024). This vector then serves as the input for a machine learning classifier. Feature fusion is particularly effective for TinyML applications, as it reduces the input dimensionality while preserving the discriminatory patterns necessary for accurate classification (Makapati et al., 2022).

3.3.0.1.2 Decision-Level Fusion

Decision-level fusion involves combining the independent outputs of multiple classifiers or heuristic rules. In the autonomous node, this manifests as the integration of TinyML model probabilities with hard threshold checks from specific “truth sensors.” For example, a “fire” classification from the neural network may only trigger a high-confidence alarm if it is confirmed by a secondary heuristic, such as a localized CO spike or a persistent flame flicker signal (Liu et al., 2023). This hybrid approach enhances the system’s robustness against transient sensor noise and isolated model errors.

3.3.0.2 Redundant and Complementary Fusion

The effectiveness of fusion in the fire detection domain relies on both redundant and complementary sensor configurations.

- **Redundant Fusion:** Multiple sensors of the same or similar types (e.g., Smoke and VOC sensors) monitor the same phenomenon. This provides fault tolerance and improves the signal-to-noise ratio, as a genuine fire will typically influence multiple sensors simultaneously (Meleti & Razi, 2024).
- **Complementary Fusion:** Sensors monitor different, yet related, physical phenomena (e.g., gas concentration vs. IR radiation). Complementary fusion is critical for false alarm rejection. As identified in the data analysis, the CO sensor acts as a “truth sensor” because its readings are rarely elevated in non-combustion scenarios like cooking steam, effectively contextualizing high smoke readings (Wang & others, 2024).

3.3.0.3 Non-Linear Discrimination via Machine Learning

Traditional fusion systems often rely on simple weighted averages or static rule-based logic. However, the complex, time-varying relationships between sensor modalities during a fire event are often non-linear. Modern intelligent edge nodes utilize machine learning—specifically neural networks—to learn these high-dimensional decision boundaries. By training on diverse datasets that include “fire,” “no_fire,” and specific “false_alarm” scenarios, the model can identify subtle patterns, such as the relationship between rising temperature and stable CO levels, to distinguish a kitchen’s ambient heat from an incipient fire (Pathan & others, 2024).

3.4 Feature Engineering for Fire Detection

Feature engineering constitutes the process of transforming raw, high-frequency sensor time series into compact, discriminative representations that maximise a classifier’s ability to separate fire, no_fire, and false-alarm conditions. Because the five modalities employed in this system—smoke, VOC, CO, IR flame, and temperature/humidity—each evolve at different temporal scales and carry distinct physical information, an effective feature set must span the time domain, the frequency domain, and the inter-sensor relational space.

3.4.1 Statistical Analysis (Mean, RMS, Kurtosis) for Gas/Thermal Trends

The most direct description of a sensor window is a set of first- and higher-order statistical moments computed over the raw readings. For gas and thermal channels that evolve slowly relative to the sampling period, these time-domain summary statistics encode the magnitude, energy, and distributional shape of the signal trajectory associated with each combustion state.

The arithmetic mean captures the average concentration or temperature level within a sliding window. During fire events, CO and VOC readings rise monotonically, so the windowed mean reflects cumulative combustion product build-up. Research has shown that CO and VOC are the most discriminative early-fire indicators in gas-sensor arrays, precisely because their mean concentrations diverge from ambient baselines well before smoke particles become detectable (Fonollosa et al., 2018). During nuisance scenarios such as cooking fumes or alcohol evaporation, VOC means also rise, but CO means do not, creating a separable mean-feature vector across classes—a property this project exploits for three-class discrimination.

The root-mean-square (RMS) value integrates signal magnitude with energy content, and is particularly informative when a channel alternates between elevated and baseline readings within a window, as occurs during intermittent heating or transient cooking events. Studies have adopted energy-normalised representations for multi-sensor gas data, reporting that energy-level features substantially reduced false-alarm confusion between cooking and genuine fire states (Kim et al., 2024).

Kurtosis quantifies the “tailedness” of the sample distribution; a high kurtosis indicates sharp, impulsive deviations from the mean, characteristic of sudden ignition transients, whereas a kurtosis near 3 (mesokurtic) reflects smooth or Gaussian variation typical of ambient or cooking conditions. Empirical verification has shown that in multi-stage feature-selection using VOC, CO₂, and temperature sensors for incipient fire classification, higher-order statistical features (including kurtosis-derived variants) increased classifier accuracy beyond what was achievable with mean-only features (Zakaria et al., 2016).

3.4.2 Frequency-Domain Features (FFT/Spectral Power) for Flame Flicker

Turbulent diffusion flames exhibit a characteristic luminous flickering caused by the periodic shedding of toroidal vortices at the flame base. This phenomenon produces oscillations in both radiated infrared intensity and, indirectly, in local gas concentrations at a frequency band of approximately 1–15 Hz, with the dominant puffing frequency of small uncontrolled flames concentrated near 10–13 Hz (Fonollosa et al., 2018).

The Fast Fourier Transform (FFT) converts a windowed time-series segment into its frequency-domain representation. From the complex spectrum, the spectral power in a band of interest is obtained. For the IR flame sensor channel, energy concentrated in the 10–15 Hz band is a strong positive indicator of a genuine flame, because static heat sources (e.g., cooking hotplates) and direct sunlight do not produce periodic oscillations in this band (Fonollosa et al., 2018).

Beyond band energy, additional spectral descriptors extracted from the FFT magnitude spectrum improve discriminability, such as the spectral centroid, spectral entropy, and peak frequency. Research has demonstrated that multi-sensor fire detection systems benefit from trend- and frequency-based feature extraction, showing that temporal pattern features derived from gas sensor signals improved model accuracy by reducing confusion with non-fire thermal events (Wu et al., 2021).

3.4.3 Cross-Sensor Correlations and Dimensionality Reduction

The strongest discriminative signal in a multi-modal system often resides in the relationships between channels. During genuine combustion, CO, VOC, temperature, and IR intensity co-vary in a physically constrained manner: CO and VOC concentrations both rise because they share the same combustion source, temperature increases as a consequence of heat release, and IR intensity fluctuates at the flicker frequency (Fonollosa et al., 2018).




Figure 1: Sensor Correlation Matrix. A heatmap showing the Pearson correlation coefficients between the six sensors. Low correlation between certain sensors (e.g., CO and Humidity) indicates they provide unique, non-redundant information for the fusion model.

Pearson’s cross-correlation coefficient between channels provides a windowed coupling measure that is physically interpretable and computationally inexpensive. However, computing all pairwise correlations across multiple sensor channels yields a high-dimensional feature vector. Principal Component Analysis (PCA) addresses this by projecting the feature matrix onto its eigenvectors of maximum variance, retaining only the components needed to explain most of the total variance while discarding redundant or noisy dimensions. Studies have demonstrated this benefit directly in incipient fire detection, where PCA reduced feature dimensionality without information loss and improved classification accuracy (Zakaria et al., 2016).

Alternative supervised dimensionality reduction methods, such as Linear Discriminant Analysis (LDA), maximize between-class scatter relative to within-class scatter. These methods are essential for generalizing learned multi-sensor fire signatures across different environments (Vorwerk et al., 2024). In the context of this project, PCA is used to reduce the combined feature vector to a fixed-length input for the neural network classifier.

3.5 Neural Network Fundamentals for Classification

Artificial Neural Networks (ANNs), specifically Multi-Layer Perceptrons (MLPs), provide the mathematical framework for performing non-linear classification on multi-modal sensor data. In the context of an autonomous fire detection node, the MLP serves as the “inference engine” that maps a vector of engineered features—derived from smoke, gas, IR, and environmental sensors—to one of three discrete classes: fire, no_fire, or false_alarm. This

section outlines the fundamental components and operational principles of neural networks as applied to this classification task.

3.5.1 The Artificial Neuron and Layered Architecture

The fundamental building block of an ANN is the artificial neuron, a computational unit that models the weighted summation of inputs followed by a non-linear activation. For a given input vector x , the neuron computes an output y as:

$$y = \sigma\left(\sum_{i=1}^n w_i x_i + b\right)$$

where w_i represents the synaptic weights, b is the bias term, and σ is the activation function (Makapati et al., 2022). In an MLP, these neurons are organized into a structured hierarchy consisting of an input layer, one or more hidden layers, and an output layer. The hidden layers enable the network to learn complex, high-dimensional decision boundaries that are inaccessible to linear classifiers (Pathan & others, 2024).

3.5.2 Activation Functions: ReLU and Softmax

Activation functions introduce the non-linearity necessary for the network to approximate complex functions.

- **Rectified Linear Unit (ReLU):** In the hidden layers, the ReLU function, defined as $f(x) = \max(0, x)$, is commonly employed. ReLU is computationally efficient for TinyML applications as it involves only a simple thresholding operation, and it mitigates the vanishing gradient problem during training (Makapati et al., 2022).
- **Softmax:** For multi-class classification, the output layer typically utilizes the Softmax activation function. Softmax squashes a vector of K real values into a probability distribution consisting of K probabilities proportional to the exponentials of the input numbers. This allows the autonomous node to interpret the model’s output as the confidence or probability of each class (fire, no_fire, or false_alarm) (Liu et al., 2023).

3.5.3 Training via Backpropagation and Gradient Descent

The process of “learning” in a neural network involves adjusting the weights and biases to minimize a predefined loss function. For a three-class classification problem, the Categorical Cross-Entropy loss is typically used, which measures the dissimilarity between the predicted probability distribution and the ground-truth one-hot encoded labels (Pathan & others, 2024).

Optimization is achieved through Backpropagation, an algorithm that calculates the gradient of the loss function with respect to each weight by applying the chain rule of calculus. These gradients are then used by an optimizer, such as Adam or Stochastic Gradient Descent (SGD), to iteratively update the weights in the direction that reduces the loss (Meleti & Razi, 2024). In TinyML workflows, this training is performed on a resource-rich machine (e.g., via the Edge Impulse platform) before the optimized, frozen model is deployed for inference on the microcontroller.

3.5.4 Non-Linear Multi-Class Discrimination

The strength of the MLP in fire detection lies in its ability to identify cross-modal correlations. While a simple heuristic might fail to distinguish cooking fumes from smoke, a trained MLP can learn that a rise in VOCs accompanied by stable CO levels and the absence

of IR flicker most likely represents a false alarm (Liu et al., 2023). This capacity for non-linear discrimination is what enables the autonomous node to achieve high precision and a low false-alarm rate across diverse environmental scenarios.

3.6 Model Quantization and Optimization for Resource-Constrained Devices

The deployment of deep learning models on resource-constrained microcontrollers, such as the ARM Cortex-M4, necessitates specialized optimization techniques to overcome severe limitations in memory (RAM), storage (Flash), and computational power. Model quantization and architectural optimization are the primary strategies used to bridge the gap between high-performance neural networks and the “edge” environment. This section examines the theoretical foundations of quantization and its impact on the performance of autonomous sensing nodes.

3.6.1 Principles of Neural Network Quantization

Quantization is the process of mapping continuous, high-precision floating-point values (typically 32-bit Float) to a discrete, lower-precision set of values (e.g., 8-bit Integer). In a neural network, this transformation is applied to both the weights and the activations.

3.6.1.1 Post-Training Quantization (PTQ)

Post-Training Quantization is a common TinyML optimization technique where a pre-trained floating-point model is converted to a fixed-point representation using a representative dataset to calibrate the dynamic range of activations. For an 8-bit integer (INT8) quantization, the mapping is defined as:

$$q = \lfloor \frac{r}{S} + Z \rfloor$$

where r is the real value, S is a scaling factor, and Z is the zero-point offset (Makkapati et al., 2022). By converting weights from 4-byte floats to 1-byte integers, the model’s storage footprint is reduced by approximately 75%, allowing complex models to fit within the 256 KB Flash of the Renesas RA4M1 microcontroller used in this project.

3.6.1.2 Quantization-Aware Training (QAT)

While PTQ is efficient, it can introduce quantization noise that degrades classification accuracy. Quantization-Aware Training (QAT) addresses this by simulating the effects of quantization during the training process itself. This allows the network to adapt its weights to the lower-precision representation, often resulting in higher accuracy than PTQ, particularly for highly compressed models (Meleti & Razi, 2024).

3.6.2 Performance Optimization and Latency Reduction

Beyond memory savings, quantization significantly enhances inference speed. Microcontrollers like the Renesas RA4M1 feature specialized instructions (e.g., SIMD - Single Instruction, Multiple Data) that can process multiple integer operations in a single clock cycle. By utilizing these instructions, an INT8-quantized model can achieve a 2x to 4x reduction in inference latency compared to its floating-point counterpart, which is critical for meeting the < 100ms real-time response target for fire detection (Makkapati et al., 2022).

3.6.3 Pruning and Architectural Compression

Complementary to quantization, pruning involves removing redundant or near-zero weights from the network, effectively “thinning” the model architecture. This reduces the number of multiply-accumulate (MAC) operations required for inference, further decreasing power consumption and latency (Meleti & Razi, 2024). For autonomous nodes, these optimizations ensure that the system can maintain high-frequency sampling (10 Hz) while simultaneously performing local inference without exhausting the device’s resources.

3.7 Overview of the Edge Impulse Platform Pipeline

The development and deployment of the machine learning model for the autonomous fire detection node are facilitated by the Edge Impulse platform, an integrated development environment specifically designed for TinyML. Edge Impulse provides an end-to-end pipeline that streamlines the transition from raw sensor data to optimized, on-device inference. This section provides an overview of the key stages in the Edge Impulse pipeline as implemented in this research.

3.7.1 Data Ingestion and Labeling

The first stage of the pipeline involves the ingestion of high-frequency sensor data. In this project, data from the multi-sensor array is streamed from the Arduino UNO R4 WiFi to the Edge Impulse studio via the edge-impulse-data-forwarder or as pre-recorded CSV files (Edge Impulse, 2024). During this phase, data is organized into three distinct classes—fire, no_fire, and false_alarm—forming the basis for supervised learning. The platform’s interface allows for precise windowing and cropping of the data to ensure that only the most relevant signal segments are used for training.

3.7.2 Impulse Design and Digital Signal Processing (DSP)

The “Impulse” represents the complete data processing block within the platform. It consists of three primary components:

1. **Input Block:** Defines the time-series window size (e.g., 10 seconds) and the sampling frequency (10 Hz).
2. **DSP Block:** Performs feature extraction to reduce the dimensionality and complexity of the raw data. In this research, a “Spectral Analysis” block is typically employed to extract both time-domain (mean, RMS) and frequency-domain (FFT) features, which are critical for identifying flame flicker and gas trends (Fonollosa et al., 2018).
3. **Learning Block:** Specifies the neural network architecture, such as a Multi-Layer Perceptron (MLP), for classification.

3.7.3 Model Training and Validation

Once the features are extracted, the platform facilitates the training of the neural network. Edge Impulse provides a cloud-based environment for hyperparameter tuning—allowing for the adjustment of layer depth, neuron count, and learning rates—and offers real-time visualization of training and validation metrics (accuracy and loss). The platform’s “Data Explorer” tool is particularly useful for visualizing the separability of classes in the feature space using dimensionality reduction techniques like T-distributed Stochastic Neighbor Embedding (t-SNE) (Edge Impulse, 2024).

3.7.4 Deployment and Optimization

The final stage of the pipeline is the deployment of the trained model to the microcontroller. Edge Impulse utilizes the EON Compiler to optimize the model for specific hardware architectures, such as the ARM Cortex-M4. This optimization includes the application of Post-Training Quantization (PTQ) to convert the model to an 8-bit integer (INT8) representation, significantly reducing RAM and Flash usage while maintaining inference speed (Makkapati et al., 2022). The model is finally exported as a self-contained C++ library, which is integrated into the Arduino firmware for real-time, on-device fire detection.

4 Sensor Selection and Characterization

4.1 Requirements for Multi-Modal Sensing

The design of an autonomous fire detection node requires a multi-modal sensing strategy to overcome the physical and environmental limitations of traditional single-parameter systems. Because fire is a complex phenomenon characterized by the simultaneous evolution of particulates, gases, infrared radiation, and heat, a single sensor modality is insufficient for robust discrimination between genuine fires and non-fire nuisance events. This section defines the technical and operational requirements that guided the selection of the multi-modal sensor suite for this research.

4.1.1 Functional and Physical Requirements

The primary functional requirement for the sensor suite is the ability to provide a comprehensive and redundant “fingerprint” of the environment. This necessitates the integration of heterogeneous sensors that target diverse physical properties:

1. **Chemical Markers:** Continuous monitoring of combustion products, specifically smoke particulates, Volatile Organic Compounds (VOCs), and Carbon Monoxide (CO), is essential for incipient fire detection (Fonollosa et al., 2018).
2. **Optical Signature:** The detection of active flames via infrared (IR) radiation provides a high-confidence indicator of open combustion, distinguishing it from smoldering events.
3. **Environmental Context:** High-resolution temperature and humidity data are required to contextualize gas sensor readings and identify thermal anomalies (Liu et al., 2023).

Physically, the sensors must be suitable for edge deployment in an “Autonomous Node” architecture. This requires low power consumption—allowing for long-term operation on battery or energy-harvesting systems—and a miniature footprint (MEMS-based) to ensure a compact and unobtrusive device design (Meleti & Razi, 2024).

4.1.2 Operational and Integration Requirements

From an operational perspective, the sensors must support high-frequency sampling (at least 10 Hz) to capture the rapid transients and temporal modulation (e.g., flame flicker) characteristic of fire events. Integration requirements include compatibility with the Arduino UNO R4 WiFi platform, necessitating sensors with standard I2C or analog interfaces. Furthermore, the gas sensors must exhibit sufficient sensitivity to detect sub-ppm concentrations of combustion products while maintaining long-term stability in diverse indoor environments (Wang & others, 2024).

4.2 Detailed Examination of Selected Sensors

This section provides a detailed examination of the hardware components selected for the autonomous fire detection node. The sensor suite comprises five distinct modalities, each chosen for its specific sensitivity to combustion products or environmental variables. All gas sensors belong to the DFRobot Fermion MEMS series, optimized for low power consumption and high integration in edge-sensing applications.

Sensor Model	Modality	Detection Range	Interface	Key Advantage
SEN0570	Smoke/ Particulate	Analog (Rel.)	Analog	Ethanol Compatible (Low False Positive)
DFR0076	IR Flame	760nm – 1100nm	Analog	< 1ms Response Time (Flicker Analysis)
SEN0566	VOC Gas	0 – 1000 ppm	Analog	Rapid Response to Chemical Vapors
SEN0564	Carbon Monoxide	1 – 1000 ppm	Analog	High Selectivity (Combustion Truth)
AHT20	Temp/Humidity	−40 to +85°C	I2C	±0.3°C Accuracy (Heat Paradox Analysis)

Table 2: Summary of Selected Multi-Sensor Suite Specifications

4.2.0.1 Fermion MEMS Smoke Detection Sensor (SEN0570)

The SEN0570 is a Micro-Electro-Mechanical Systems (MEMS) based smoke sensor designed for high-sensitivity detection of particulates associated with combustion. Unlike traditional ionization or photoelectric detectors, this MEMS implementation utilizes a metal-oxide semiconductor (MOS) layer that reacts to a wide range of smoke types, including those from wood, paper, and synthetic materials. A significant operational advantage of the SEN0570 is its ethanol compatibility, which allows it to maintain stability in environments where alcohol-based cleaning products might trigger false positives in less sophisticated sensors (DFRobot, 2024a). The sensor operates on a 3.3V–5V supply and provides an analog output proportional to smoke concentration, making it ideal for integration with the high-resolution ADCs of the Renesas RA4M1 microcontroller.

4.2.0.2 Gravity Analog Flame Sensor (DFR0076)

The Gravity Analog Flame Sensor is an infrared (IR) based device optimized for detecting radiation in the 760 nm to 1100 nm wavelength band. This spectrum is characteristic of the thermal emissions of an open flame. The sensor utilizes a phototransistor with a wide detection angle of approximately 60 degrees and a response time of less than 1 ms, enabling the capture of high-frequency flame flicker (Toreyin et al., 2012). The onboard potentiometer

allows for hardware-level sensitivity adjustment, providing a first line of defense against static IR interference before digital signal processing.

4.2.0.3 Fermion VOC Gas Sensor (SEN0566)

Volatile Organic Compounds (VOCs) are primary indicators of early-stage fire, especially when synthetic materials or chemical accelerants are involved. The SEN0566 MEMS VOC sensor is designed to detect a broad spectrum of organic vapors, including formaldehyde, benzene, and toluene, with a detection range of 0 to 1000 ppm (DFRobot, 2024b). Its small form factor and low power consumption (< 50 mW) are critical for continuous 24/7 monitoring in stationary nodes. The sensor’s rapid response to ambient air quality changes provides the machine learning model with the high-frequency temporal data necessary to distinguish between slow-building fire signatures and transient environmental changes.

4.2.0.4 Fermion MEMS Carbon Monoxide (CO) Sensor (SEN0564)

The detection of Carbon Monoxide (CO) is the most reliable method for combustion verification, as CO is a byproduct of nearly all fire events but is rarely produced by common household nuisance sources like steam or aerosol sprays. The SEN0564 MEMS CO sensor provides high selectivity for CO molecules with a typical sensing range of 1 to 1000 ppm (DFRobot, 2024c). By acting as a “truth sensor” within the fusion model, the CO sensor significantly reduces the false alarm rate by providing a physical confirmation of incomplete combustion that must coincide with elevated smoke or VOC readings (Liu et al., 2023).

4.2.0.5 Fermion AHT20 Temperature and Humidity Sensor

Environmental context is provided by the AHT20, a high-precision digital sensor that communicates via the I2C protocol. It offers a temperature accuracy of $\pm 0.3^{\circ}\text{C}$ and a relative humidity accuracy of $\pm 2\%$. In this research, the AHT20 serves two critical roles: first, it provides data for the “heat paradox” analysis, where thermal anomalies are used to distinguish cooking events from incipient fires (Meleti & Razi, 2024); second, it enables humidity-based compensation for the MOS gas sensors, whose resistance can be influenced by ambient moisture levels (Pathan & others, 2024).

4.3 Justification of Sensor Choice

The selection of the multi-modal sensor suite for the autonomous fire detection node is justified by the need to balance high sensitivity for early-stage fire detection with high selectivity to minimize false alarms. Traditional fire detectors often suffer from a trade-off where increasing sensitivity leads to a proportional increase in nuisance alarms from non-fire aerosols or environmental fluctuations. This section provides the rationale for the chosen sensors, focusing on how their complementary properties address this fundamental challenge.

4.3.1 Sensitivity for Early Detection

The primary justification for selecting the DFRobot Fermion MEMS gas sensor series (Smoke and VOC) is their exceptional sensitivity to the gaseous and particulate byproducts of incipient fires. MEMS-based metal-oxide (MOS) sensors exhibit rapid response times and can detect combustion markers at sub-ppm concentrations, often before visible smoke or significant heat is present (Fonollosa et al., 2018). By integrating both a broad-spectrum VOC sensor and a particulate-focused smoke sensor, the system captures a high-resolution

“chemical snapshot” of the environment, ensuring that slow-building or smoldering fires are identified in their earliest stages.

4.3.2 Selectivity and the “Truth Sensor” Strategy

While sensitivity ensures detection, selectivity is critical for discrimination. The inclusion of the SEN0564 MEMS Carbon Monoxide (CO) sensor is a strategic choice justified by its high selectivity for CO, a universal product of incomplete combustion. As demonstrated in the project’s data analysis, CO levels remain near baseline in common false alarm scenarios such as cooking steam, aerosol sprays, or cleaning alcohol, whereas they spike significantly during genuine fire events (Liu et al., 2023). By designating the CO sensor as a “truth sensor” within the fusion model, the node can effectively “veto” high readings from the more cross-sensitive smoke and VOC sensors if a corresponding CO rise is not detected, significantly reducing the false positive rate.

4.3.3 Redundancy and Contextual Validation

The Gravity Analog Flame Sensor and the AHT20 environmental sensor provide the necessary redundant and contextual data to validate chemical sensor readings. The IR flame sensor offers a non-chemical detection modality that is nearly instantaneous, providing a high-confidence indicator of open flames that is physically distinct from gas-based detection (Toreyin et al., 2012). Simultaneously, the AHT20’s temperature and humidity data are used to contextualize gas sensor resistance changes, allowing the machine learning model to account for ambient environmental variance and identify “heat paradox” scenarios where false alarms may exhibit higher localized temperatures than early-stage fires (Meleti & Razi, 2024). This multi-layered approach ensures that the node’s final classification is based on a fused consensus of physically diverse phenomena.

4.4 Sensor Calibration Protocols and Baseline Behavior

The reliability of an autonomous fire detection node is contingent upon the accurate calibration of its sensors and a clear understanding of their baseline behavior in diverse environments. Because metal-oxide (MOS) gas sensors and infrared flame sensors are prone to drift and environmental interference, specific protocols must be implemented during both hardware initialization and software execution to ensure data integrity. This section outlines the calibration procedures and baseline stabilization techniques employed in the project’s firmware.

4.4.1 Hardware Pre-heating and Stabilization

MOS-based MEMS gas sensors, such as the Fermion smoke, VOC, and CO sensors, require a thermal stabilization period to achieve a consistent baseline resistance. Upon initial deployment, new sensors must undergo a “burn-in” period of approximately 48 hours to remove contaminants from the sensing layer and stabilize the internal micro-heater (DFRobot, 2024b). In daily operation, the node implements a software-controlled warm-up phase of 120 seconds upon power-up, during which sensor readings are monitored but not used for inference. This ensures that the sensing layer has reached its optimal operating temperature (typically 200°C to 400°C), minimizing the risk of false positives caused by cold-start transients (Pathan & others, 2024).

4.4.2 Software-Level Baseline Detection

To account for long-term sensor drift and varying ambient air quality, the system utilizes a dynamic baseline detection algorithm within the Arduino firmware. The baseline represents the “no_fire” state of the environment.

4.4.2.1 Differential Sensing Logic

Rather than relying on absolute voltage thresholds, the node’s inference logic is based on the differential change from the established baseline ($\Delta R/R_0$). During the first 10 seconds of stable operation (after the warm-up period), the firmware calculates the mean and standard deviation of each gas sensor channel to define the ambient baseline (Wang & others, 2024). This allows the system to adapt to different rooms or seasonal variations in background gas concentrations. As implemented in the main firmware, the system performs continuous “sanity checks” to ensure that the baseline remains within expected physiological limits before permitting the machine learning model to trigger an alarm.

4.4.3 Flame Sensor Sensitivity Adjustment

The Gravity Analog Flame Sensor requires a manual calibration of its onboard potentiometer to define the detection range. Calibration is performed by exposing the sensor to a controlled IR source (simulating a small flame) at the maximum required detection distance and adjusting the potentiometer until the digital trigger threshold is reached. This hardware-level adjustment provides a coarse filter for static IR noise, such as sunlight or high-intensity artificial lighting, which is then further refined by the temporal flicker analysis in the machine learning model (Meleti & Razi, 2024).

4.4.4 Environmental Compensation (AHT20)

The AHT20 sensor comes factory-calibrated and does not require user adjustment. However, its high-resolution temperature and humidity data are used to compensate for the cross-sensitivity of the MOS gas sensors. As humidity increases, water vapor molecules compete for active sites on the gas sensing layer, potentially lowering the perceived resistance (Pathan & others, 2024). The fusion model uses the AHT20’s digital output to normalize these readings, ensuring that the detected gas trends are independent of ambient weather conditions.

5 Hardware Platform and System Integration

5.1 The Autonomous Edge-Node Architecture

The “Autonomous Sensing Node” architecture developed in this research represents a paradigm shift from traditional, centralized fire detection systems toward intelligent, edge-deployed units capable of local decision-making. This architecture is physically and computationally centered on the Arduino UNO R4 WiFi, a versatile platform that enables the fusion of multi-modal sensor data with real-time TinyML inference. This section justifies the selection of this hardware platform and details the Asymmetric Multi-Processing (AMP) strategy that defines its operation.

5.1.1 Justification for the Arduino UNO R4 WiFi Platform

The selection of the Arduino UNO R4 WiFi as the primary controller was guided by three critical requirements: computational overhead for TinyML, multi-protocol connectivity, and ecosystem compatibility. Unlike its 8-bit predecessor (the UNO R3), the R4 WiFi is powered by the Renesas RA4M1 microcontroller, featuring a 32-bit ARM Cortex-M4 core running at 48 MHz (Arduino, 2024).

This hardware upgrade provides the necessary 256 KB of Flash and 32 KB of SRAM required to host complex neural network models generated by the Edge Impulse platform. Furthermore, the RA4M1 includes a built-in Floating Point Unit (FPU) and support for DSP instructions (CMSIS-DSP), which are essential for high-frequency feature extraction, such as the 128-point FFTs used for flame flicker analysis (Makkapati et al., 2022). The integration of an ESP32-S3 module alongside the RA4M1 provides a seamless bridge to modern IoT protocols without burdening the primary inference core, satisfying the requirement for real-time MQTT telemetry in stationary nodes.

5.1.2 Asymmetric Multi-Processing (AMP) Strategy

The core innovation of the R4 WiFi architecture is its Asymmetric Multi-Processing (AMP) capability, which allows for the functional separation of sensing, inference, and communication tasks across two distinct processing units.

Feature	RA4M1 “Inference Brain”	ESP32-S3 “Connectivity Backbone”
Primary Role	Sensing, DSP, TinyML Inference	WiFi, MQTT, Security, Telemetry
Core Architecture	ARM Cortex-M4 (32-bit)	Xtensa LX7 (Dual-core)
Clock Speed	48 MHz	240 MHz
Critical Resource	Floating Point Unit (FPU)	Integrated WiFi/Bluetooth Radio
Data Flow	Master (Decision Maker)	Slave (Communication Bridge)

Table 3: Functional Separation in the Arduino UNO R4 WiFi AMP Architecture

5.1.2.1 The RA4M1 as the “Inference Brain”

In this node architecture, the Renesas RA4M1 serves as the “Inference Brain.” It is responsible for the deterministic, time-critical tasks of the fire detection pipeline:

1. **Sensor Data Acquisition:** High-frequency sampling (10 Hz) of the five heterogeneous sensor modalities via I2C and high-resolution (14-bit) analog inputs.
2. **Digital Signal Processing (DSP):** Local computation of mean, RMS, and spectral power features within sliding windows.
3. **TinyML Inference:** Execution of the quantized INT8 neural network model to classify the environment into fire, no_fire, or false_alarm states.

By isolating these tasks on the RA4M1, the system ensures that the sensing and inference loop is never interrupted by network latency or connectivity overhead (Meleti & Razi, 2024).

5.1.2.2 The ESP32-S3 as the “Connectivity Backbone”

The ESP32-S3 serves as the “Connectivity Backbone,” acting as a specialized co-processor for wireless communication. It handles the high-level networking stack, including WiFi association, security protocols, and MQTT message serialization. When the RA4M1 identifies a significant event or reaches a classification consensus, it communicates the metadata to the ESP32-S3 via a dedicated serial bridge (UART). This modular approach allows the ESP32-S3 to manage telemetry and remote logging in the background, ensuring that the autonomous node remains responsive even during periods of heavy network traffic (Deng & others, 2023). This functional split between the “brain” and the “backbone” is critical for the reliability of life-safety systems, as it prevents communication failures from compromising local detection capabilities.

5.2 Power Management and Thermal Considerations

The power management and thermal design of the autonomous sensing node are critical factors that directly influence sensor stability and detection accuracy. Unlike low-power environmental loggers, a multi-sensor fire detection unit must support the continuous operation of active heating elements within its gas sensors, necessitating a robust power delivery system and effective heat dissipation strategies.

5.2.1 Power Demands of MEMS Gas Sensors

The primary power consumption in the autonomous node is driven by the internal micro-heaters of the three DFRobot Fermion MEMS gas sensors (Smoke, VOC, and CO). Metal-oxide (MOS) sensors require these heaters to maintain the sensing layer at high temperatures to facilitate the redox reactions necessary for gas detection (Pathan & others, 2024).

Each MEMS sensor consumes significant current depending on its specific operating state and target gas. When combined with the Arduino UNO R4 WiFi’s processors, the total peak current draw of the node can exceed 500 mA. To maintain data integrity, the node utilizes the stabilized 5V rail from the Arduino’s onboard voltage regulator. A stable supply is essential because fluctuations in the heater voltage can lead to baseline drift and increased signal noise, which can be misidentified by the TinyML model as a fire-related trend (Makkapati et al., 2022).

5.2.2 Thermal Management and Self-Heating

The concentration of three active micro-heaters within a compact enclosure presents a significant thermal challenge. This “self-heating” effect can create a localized micro-climate inside the device that is warmer than the ambient room temperature.

5.2.2.1 Impact on Environmental Sensing

If not properly managed, the heat generated by the gas sensors can interfere with the environmental sensor, leading to inaccurate context data. This is particularly problematic for the “heat paradox” analysis, where the system must distinguish between external fire-induced heat and internal device heat (Meleti & Razi, 2024). To mitigate this, the hardware integration strategy involves physically separating the environmental sensor from the gas

sensor array and utilizing a ventilated enclosure design. Ventilation ensures continuous airflow over the sensors, allowing the heaters to maintain their required internal temperatures while preventing the buildup of stagnant hot air that would skew the external temperature readings (Wang & others, 2024).

5.3 Physical Design and Enclosure

The physical design and enclosure of the autonomous fire detection node are engineered to optimize sensor exposure while maintaining the structural and thermal integrity of the electronic components. Unlike general-purpose electronics enclosures, a multi-sensor fire detection unit requires a specialized design that facilitates continuous gas exchange and provides robust isolation between internal heat sources and sensitive environmental probes.

5.3.1 Airflow Considerations for Gas Sensing

The primary requirement for the enclosure is to ensure that ambient air can freely circulate over the metal-oxide (MOS) gas sensor array. Because the detection of smoke, VOCs, and CO relies on the physical adsorption of gas molecules onto the sensing layer, any restriction in airflow can lead to significant detection delays, potentially exceeding the critical response window for incipient fires (Fonollosa et al., 2018).

The enclosure design utilizes a “ventilated chamber” approach, featuring strategically placed intake and exhaust vents that leverage natural convection. The heat generated by the sensors’ internal micro-heaters creates a localized “chimney effect,” where rising warm air draws in fresh ambient air from the lower vents. This passive airflow strategy ensures that the sensors are continuously exposed to a representative sample of the room’s atmosphere without the need for active fans, which would increase power consumption and complexity (Wang & others, 2024).

5.3.2 Thermal and Modal Isolation

A secondary but equally critical design objective is the physical and thermal isolation of the node’s constituent components.

5.3.2.1 Isolation of Environmental Sensors

As detailed in previous sections, the MOS gas sensors operate at high internal temperatures, generating localized heat that can skew environmental readings. To mitigate this, the node’s layout follows a “modal isolation” strategy. The digital temperature and humidity sensor is positioned in a thermally isolated compartment, separated from the gas sensor array by a physical barrier or a minimum clearance distance (Meleti & Razi, 2024). This ensures that the sensor captures the true ambient conditions of the room rather than the self-heating effects of the device, which is essential for accurate cross-sensitivity compensation and “heat paradox” detection.

5.3.2.2 Optical Alignment for Flame Detection

The physical design also accounts for the line-of-sight requirements of the infrared flame sensor. The enclosure features a dedicated optical aperture with a clear IR-transparent window or an unobstructed opening, aligned to provide a 60-degree field of view (Liu et al., 2023). This positioning ensures that the IR sensor can capture radiant emissions from open flames while being protected from mechanical damage and dust accumulation.

5.4 Connectivity and Location Awareness

The integration of robust connectivity and location awareness is essential for transforming a standalone fire detection node into a networked safety system. For autonomous sensing nodes, connectivity enables real-time telemetry and remote alerting, while location awareness provides the critical spatial context necessary for effective emergency response. This section details the MQTT-based communication architecture and the static location tagging strategy implemented in this research.

5.4.1 MQTT Telemetry and Network Architecture

Communication in the autonomous node is based on the Message Queuing Telemetry Transport (MQTT) protocol, a lightweight publish-subscribe messaging standard optimized for resource-constrained IoT devices. This protocol is particularly suited for fire detection applications due to its low overhead and support for varying Quality of Service (QoS) levels, ensuring that critical alarm messages are delivered reliably even in unstable network conditions.

As detailed in the system architecture, the ESP32-S3 co-processor manages the high-level networking stack, including WiFi association and MQTT message serialization. The node publishes data to a centralized MQTT broker under a structured topic hierarchy. The telemetry payload, formatted in JSON, includes the TinyML model’s classification probabilities, raw sensor snapshots during significant events, and internal device health metrics (Deng & others, 2023). This decoupled approach allows the primary RA4M1 processor to remain focused on sensing and inference, while the ESP32-S3 handles the asynchronous communication tasks in the background, minimizing the risk of network-induced latency in the detection loop (Meleti & Razi, 2024).

5.4.2 Static Location Tagging and Zone Identification

While modern mobile devices utilize GNSS or Wi-Fi fingerprinting for dynamic positioning, stationary autonomous nodes rely on a static “Zone_ID” implementation for location awareness. In the context of building fire safety, identifying the specific room or floor of an incipient fire is more critical than high-precision coordinates.

Each node is assigned a unique Zone_ID during the firmware configuration phase, which is stored in the RA4M1’s non-volatile memory. This identifier corresponds to a physical location within the building’s digital twin or floor plan (e.g., “Kitchen_A1”). When a “fire” state is identified by the on-device model, the classification metadata published via MQTT is automatically enriched with this Zone_ID. This spatial context allows building management systems and emergency responders to immediately pinpoint the origin of the threat, facilitating targeted evacuation and faster suppression (Liu et al., 2023). Furthermore, the Zone_ID enables “location-aware filtering” at the broker level, where alert priorities can be dynamically adjusted based on the occupancy or hazard level of the specific zone.

5.5 System Schematic and Wiring Diagrams

The hardware integration of the autonomous sensing node is defined by a structured interconnectivity map that ensures signal integrity and reliable power delivery. The system utilizes a combination of digital and analog interfaces to accommodate the heterogeneous

sensor suite, with the Arduino UNO R4 WiFi acting as the central nexus. This section details the system’s schematic logic, focusing on the I2C bus architecture and analog pin mapping.

5.5.1 Digital Interconnectivity: The I2C Bus

Digital communication is centralized on the Inter-Integrated Circuit (I2C) bus, utilizing the SDA (Serial Data) and SCL (Serial Clock) pins of the RA4M1 microcontroller. This bus follows a multi-slave architecture, allowing for high-resolution data acquisition with minimal wiring complexity. The AHT20 temperature and humidity sensor is the primary digital slave on this bus, providing factory-calibrated environmental data at 10 Hz (Pathan & others, 2024). To ensure bus stability, pull-up resistors are integrated into the circuit, and the firmware implements a stabilization delay after initialization.

5.5.2 Analog Interface and Pin Mapping

The MOS-based gas sensors and the IR flame sensor utilize the RA4M1’s high-resolution Analog-to-Digital Converters (ADCs) for signal transduction. This mapping is critical for capturing the subtle voltage fluctuations associated with incipient fire signatures.

- **Smoke and VOC Sensors:** The MEMS smoke and VOC sensors are mapped to dedicated analog input pins. The firmware reads the voltage divider output from these sensors, where the resistance change of the sensing layer is converted into a proportional voltage signal (Wang & others, 2024).
- **Carbon Monoxide (CO) Sensor:** The MEMS CO sensor is mapped to a high-priority analog pin, reflecting its role as the “truth sensor.”
- **Infrared Flame Sensor:** The Gravity Analog Flame Sensor provides a continuous voltage output to a dedicated analog pin, enabling the capture of high-frequency flicker (Meleti & Razi, 2024).

5.5.3 Power Distribution and Grounding

The schematic implements a parallel power distribution strategy. All sensor heaters are tied to the Arduino’s 5V rail, while the digital logic of the environmental sensor is powered by the 3.3V rail. A “common ground” star topology is employed to minimize ground loops and signal interference, ensuring that the analog baselines remain stable across all modalities (Makkapati et al., 2022).

6 Data Collection Methodology

6.1 Safety Protocols During Data Collection

Conducting laboratory experiments involving controlled combustion events introduces a range of physical, chemical, and environmental hazards that demand structured risk mitigation prior to data acquisition. The generation of open flames, smoldering materials, and the associated release of toxic combustion by-products — including carbon monoxide (CO), volatile organic compounds (VOCs), and particulate matter — create conditions that can pose immediate harm to researchers if left unmanaged. Accordingly, the data collection phase of this project was preceded by a formal risk assessment and the establishment of layered safety controls spanning fire containment, laboratory ventilation, and the use of personal protective equipment (PPE).

6.1.1 Hazard Identification and Risk Assessment

Before any combustion experiment was initiated, a systematic hazard identification process was conducted to catalogue the primary risks associated with each planned fire scenario. Three categories of hazard were identified: (1) uncontrolled fire propagation, (2) inhalation of toxic combustion gases, and (3) thermal burns to personnel.

The toxic hazard profile of combustion events is well established in the literature. CO, hydrogen cyanide (HCN), and VOCs are identified as the principal toxic agents released during both smoldering and flaming fires (Fonollosa et al., 2018). Specifically, CO — produced under lean combustion conditions — can reach lethal concentrations in a confined space, and its toxic potency is further amplified when co-present with CO₂. Similarly, VOCs generated during the pyrolysis phase of polymeric materials — including acrolein and formaldehyde — represent significant irritant hazards (Fonollosa et al., 2018). Research corroborates these findings, confirming that CO remains the most diagnostic and most hazardous combustion gas across diverse fire types (Khan et al., 2022).

For this project, fire loads included paper, wood, and cloth (fire class), as well as cooking-derived aerosols and alcohol vapors (false alarm class). A risk matrix was constructed prior to experiments, mapping each scenario to its expected gas concentrations and thermal hazard level. Scenarios involving direct ignition of solid materials (e.g., wood, cloth) were classified as highest risk and subjected to the most stringent containment and monitoring controls.

6.1.2 Fire Containment Measures

All controlled burns were conducted within a designated fireproof experimental enclosure lined with non-combustible ceramic tile and stainless-steel trays to contain any ash, embers, or liquid accelerants. The combustion test volume was dimensionally bounded to prevent flame propagation beyond the immediate sensor proximity zone, consistent with small-scale fire test methodologies reported in peer-reviewed literature.

Research describes a directly analogous small-scale experimental setup for multi-sensor fire data collection, ensuring that combustion remained confined to a precisely defined fuel mass without transitioning to uncontrolled open flame (Vorwerk et al., 2024). In the present project, a dry-powder fire extinguisher was positioned within immediate reach at all times during experiments involving open flame scenarios. No experiment was conducted without a second researcher present to act as a dedicated safety monitor.

6.1.3 Ventilation Requirements

Adequate ventilation was identified as the primary engineering control for managing the inhalation hazard posed by CO and VOC accumulation during experiments. All combustion experiments were conducted in a room equipped with a mechanical exhaust ventilation system, with the exhaust outlet positioned directly above the combustion zone to capture buoyant combustion products before dispersal into the researcher’s breathing zone.

The criticality of ventilation control in fire detection experiments is underscored by research showing that experimental conditions directly affect the rate of CO, smoke, and temperature accumulation in the test room (Li et al., 2022). Standardised ventilation states were maintained for all fire and false-alarm class recordings, while a sealed-room baseline was recorded for no-fire ambient conditions to characterise background sensor drift. Following

each experiment, a mandatory purge period of no less than 10 minutes was observed before the next trial commenced, allowing sensor baselines to recover and residual gases to clear.

6.1.4 Personal Protective Equipment (PPE)

Researchers conducting experiments were required to wear a minimum standard of PPE throughout all fire class and false-alarm class data collection sessions. The mandatory PPE set comprised a half-face elastomeric respirator fitted with combined P3/A1 filter cartridges, safety glasses, a flame-resistant cotton laboratory coat, and heat-resistant silicone gloves when handling materials post-experiment.

The necessity of respiratory protection during fire-related experimental work is corroborated by research noting that experimental procedures for incipient fire generation expose personnel to sub-threshold but cumulatively significant concentrations of CO and VOC (Vorwerk et al., 2024). A personal clip-on electrochemical CO monitor was worn by the lead researcher at all times during combustion experiments. If the personal monitor alarmed, the experiment was immediately suspended, the room evacuated, and mechanical ventilation maximised. No experiment proceeded if the pre-session ambient CO reading exceeded 5 ppm, providing a conservative pre-screening margin.

6.2 Rationale for Three-Class Classification

The decision to frame fire detection as a three-class classification problem—distinguishing between fire, no_fire, and false_alarm—is a deliberate design choice motivated by the demonstrable inadequacy of binary approaches in the presence of real-world nuisance sources. This section articulates the theoretical and empirical grounds for this classification schema, drawing on the literature on false alarm prevalence, the chemical similarity between nuisance and fire events, and the comparative performance of binary versus multi-class models in sensor-based fire detection systems.

6.2.1 The Limitations of Binary Classification in Practical Fire Detection

The dominant paradigm in early sensor-based fire detection has historically been binary classification: a system either declares a fire state or a no-fire state. While conceptually straightforward, this formulation conflates two phenomenologically distinct non-fire conditions—genuine ambient environments and nuisance events—under a single negative class. Research has observed that gas-sensor arrays trained on binary fire/no-fire models still produced a high false alarm rate when exposed to nuisance scenarios not seen during training (Fonollosa et al., 2018). This finding underscores a structural weakness of the binary model: since nuisance events can produce sensor signatures that partially overlap with incipient fire signatures, collapsing both ambient and nuisance conditions into a single no-fire class forces the decision boundary to occupy an ambiguous and poorly defined feature subspace.

The prevalence of false alarms carries significant real-world consequences that further justify a more granular classification approach. False alarms from automatic fire detection systems are estimated to cost approximately £1 billion per year in certain regions (Fonollosa et al., 2018). In such an environment, a binary classifier that cannot distinguish between a cooking-induced smoke plume and actual combustion is operationally unreliable; occupants habituate to frequent spurious alarms and may disable or ignore detection systems altogether.

6.2.2 Chemical and Thermal Overlap Between Nuisance and Fire Events

The core technical challenge motivating the three-class design is the degree to which common false alarm sources produce sensor responses that mimic early fire conditions. Cooking fumes, for instance, generate aerosols and elevated temperatures that activate optical smoke sensors and thermal detectors in ways that are nearly indistinguishable from smoldering fires at the single-sensor level. Research has demonstrated this overlap empirically in a gas sensor array study, where even well-calibrated multivariate models struggled to maintain separation between fire and nuisance when the underlying chemical signatures shared common compounds (Solórzano & others, 2021).

Similar overlap has been documented at the algorithmic level. Research has reported that in a multi-class fire perception system, the most common misclassification involved confusing smoldering with no-fire observations, precisely because early-stage smoldering produces gradual, low-amplitude sensor trends that are difficult to distinguish from ambient drift (Li et al., 2022). This characteristic makes smoldering events the most dangerous scenario for a binary system, since the fire state is most likely to be suppressed under the dominant no-fire class until combustion products reach hazardous concentrations.

6.2.3 Three-Class Classification as a Design Solution

Introducing an explicit `false_alarm` class directly addresses the shortcomings described above by providing the model with labeled training examples of nuisance events—cooking fumes, alcohol vapors, steam, and intense infrared illumination—that share partial sensor signatures with fire but lack the complete multivariate profile of genuine combustion. This expands the decision space such that the classifier must learn a distinct, dedicated region for nuisance phenomena rather than allowing those observations to contaminate either the fire or the `no_fire` boundary.

The effectiveness of moving beyond binary classification has been corroborated in the literature. Research has demonstrated that classifying sensor array data into three conditions using a neural network achieved a high classification rate, establishing that the nuisance class carries discriminative information that improves overall system performance (Vorwerk et al., 2024).

For the present system, the three classes are operationally defined as follows. The fire class encompasses genuine combustion events producing correlated elevations across CO, smoke, VOC, and IR flame channels, accompanied by temperature rise. The `no_fire` class represents stable ambient conditions with no significant cross-sensor activity. The `false_alarm` class captures nuisance events that may activate one or two individual sensors but do not produce the correlated multi-sensor signature characteristic of actual fire. This tripartite schema leverages the full dimensionality of the five-sensor fusion architecture and provides the TinyML classifier with the inter-class boundaries necessary to suppress spurious activations without reducing sensitivity to genuine fire events.

6.3 Sampling Parameters

The selection of appropriate sampling parameters — encompassing both the data acquisition rate and the temporal window over which observations are aggregated — is a foundational decision in any embedded time-series classification system. In the context of multi-sensor fire detection, these parameters govern whether the node captures the transient chemical

signatures of combustion, the characteristic flicker dynamics of an open flame, and the slower thermal and gas-concentration drifts that distinguish genuine fire from nuisance events. Critically, the chosen parameters must simultaneously satisfy the physical requirements of each sensing modality and remain within the computational budget of a Cortex-M4 microcontroller running a quantized neural network in real time.

6.3.1 Sampling Frequency

The Nyquist–Shannon sampling theorem states that a signal must be sampled at a rate strictly greater than twice the highest frequency component of interest in order to avoid aliasing artefacts during reconstruction and feature extraction (Samanta & others, 2024). In the context of infrared flame sensing, this constraint has concrete physical significance: turbulent, uncontrolled flames exhibit a wide-band flicker spectrum spanning 1–13 Hz, with the characteristic central frequency of approximately 10 Hz reported consistently in the combustion literature (Toreyin et al., 2012). To satisfy the Nyquist condition for content up to the 13 Hz upper bound, a minimum acquisition rate of 26 Hz would be theoretically required for the IR channel alone; however, the present system adopts a unified, cross-modal sampling rate of 10 Hz (i.e., one sample vector every 100 ms) as a deliberate engineering trade-off.

This decision is justified by the response characteristics of the chemical sensing modalities. MEMS-based smoke, VOC, and CO sensors operate via transduction mechanisms whose physical response times typically range from several seconds to tens of seconds (Vorwerk et al., 2024). Sampling these channels at rates substantially exceeding 10 Hz yields no additional information content, as successive readings are highly auto-correlated and dominated by sensor time-constants rather than by underlying environmental dynamics. The 10 Hz rate thus represents the practical upper bound imposed by the chemical sensor ensemble, ensuring that all channels within the fused input vector are acquired at a consistent, coherent cadence.

From a system-resource perspective, a 10 Hz sampling rate also aligns with the computational throughput of the Renesas RA4M1 (Cortex-M4 at 48 MHz) running the Edge Impulse inference engine. Research has demonstrated that reducing data acquisition rates on TinyML platforms can decrease RAM usage and inference latency by approximately 60–74% with minimal classification accuracy loss (Samanta & others, 2024). The 10 Hz rate provides a 100 ms inter-sample interval that is comfortably shorter than the sub-second transient signatures associated with flaming combustion onset, while remaining well within the real-time budget required for continuous, non-blocking inference on the target hardware.

6.3.2 Window Size and Overlap Strategy

Once the per-sample cadence is established, the second critical parameter is the temporal window: the contiguous block of consecutive samples that is assembled into a single labeled instance and presented to the digital signal processing (DSP) block and subsequent classifier.

A 2-second window (20 samples at 10 Hz) was selected as the primary inference window for this system. This duration was chosen to capture sufficient temporal context to extract both time-domain statistical features (e.g., mean, root mean square, kurtosis) and frequency-domain features (e.g., spectral power in the 1–10 Hz flame-flicker band) from the IR channel with adequate frequency resolution. Research employed a 10-second window at a 10 Hz rate for an embedded fire detection CNN, demonstrating that windows of this order yield stable classification (Deng & others, 2023). A shorter 2-second window reduces the time-to-first-

detection latency — a critical factor in life-safety applications — while still encoding the dominant sub-second chemical transients and flame-flicker energy.

Second, the window length directly determines the RAM footprint of the inference buffer on the RA4M1. At 10 Hz, five sensor channels, and single-precision representation, a 2-second window occupies a negligible fraction of the 128 KB SRAM available on the microcontroller, leaving ample headroom for the model weights, activation buffers, and firmware stack.

A 50% overlap (i.e., a new inference is triggered every 1 second, with the preceding 1-second block retained) is employed to increase temporal coverage and reduce the expected worst-case detection latency. With 50% overlap, a fire event that begins mid-window will be fully enclosed within the subsequent overlapping window, capping the maximum detection latency at one window stride (1 second) rather than one full window length (2 seconds). This sliding-window approach is standard practice in embedded TinyML classification pipelines (Samanta & others, 2024).

For data collection, a capture duration of 10 seconds was used to provide sufficient samples for off-line model training and validation, ensuring that the resulting labeled training instances carry rich gas-concentration trend information — particularly important for distinguishing the gradual CO rise of smoldering combustion from the abrupt VOC spike of a cooking false-alarm event (Vorwerk et al., 2024).

6.4 Class-Specific Data Collection Scenarios

To develop a robust multi-sensor fire detection system capable of distinguishing between actual fire events, normal ambient conditions, and common false alarms, a meticulously planned data collection methodology was employed. This approach focused on generating a comprehensive dataset encompassing three distinct classes: “Fire,” “No-Fire,” and “False Alarm.” Each class was characterized by specific environmental scenarios designed to capture a wide range of relevant sensor signatures.

Class	ID	Scenario Name	Primary Physical Signature
Fire	A1	Close Range Flame	High Flame (>800), Rising Smoke/CO
Fire	A3	Smoldering Material	High Smoke/VOC, Low Flame IR
No-Fire	B1	Ambient Office	Stable Baseline (Low readings across all)
False Alarm	C1	Cooking Fumes	High VOC/Smoke, Negligible CO
False Alarm	C2	Steam / Humidity	High Humidity (>80%), Stable CO/VOC
False Alarm	C3	Alcohol Sprays	Extreme VOC Spikes, Stable Flame/CO

Table 4: Catalog of Data Collection Scenarios and Physical Signatures

The data collection adhered to strict protocols to ensure consistency, reproducibility, and safety, utilizing the automated data collection scripts (the automated data collection utility) for standardized sampling.

6.4.0.1 Fire Class: Paper, Wood, Cloth Burns

The “Fire” class aimed to capture the multi-sensor signatures of genuine combustion events under varying conditions. This involved controlled burns of common household materials to simulate realistic fire scenarios. The goal was to generate data where flame, smoke, CO, VOC, and temperature/humidity sensors would exhibit characteristic responses indicative of active fire.

- **Scenario A1: Close Range & Low Ventilation:** This scenario simulated an early-stage fire in a confined space. A small flame source (e.g., candle or gas burner) was positioned 10-15 cm directly in front of the sensor array in a closed environment. Data collection commenced after 30 seconds of flame stabilization. Verification involved checking flame values (700-900) and rising smoke (the project data collection guide, n.d.).
- **Scenario A2: Medium Range & Normal Ventilation:** This scenario represented a more developed fire with typical airflow. The flame source was placed 30-50 cm from the sensors in an open-ventilated space. flame values were expected to be moderate (400-700), reflecting the increased distance and dispersion (the project data collection guide, n.d.).




Figure 2: Indoor Fire Data Collection. Left: capturing infrared signatures from a flame source; Right: capturing combined flame and smoke signatures.

- **Scenario A3: Smoldering:** To capture the signatures of slow, smoldering fires, materials like incense sticks or extinguished matches were used. These events primarily generate smoke and VOCs without an open flame, posing a distinct detection challenge for IR-based sensors (the project data collection guide, n.d.).




Figure 3: Smoldering Fire Simulation. The sensor node capturing particulate and VOC signatures from a smoldering source without visible flame.

Each fire scenario was conducted to capture multi-sensor data at a sampling rate of 10 Hz over a specific duration, ensuring a rich temporal dataset for model training. The raw data for the “Fire” class is organized under the `fire dataset` directory.

6.4.0.2 No-Fire Class: Idle Office, Kitchen Ambient

The “No-Fire” class aimed to establish a robust baseline of normal environmental conditions. This data is critical for training the model to recognize the absence of fire and differentiate it from both true fire events and false alarms. The scenarios focused on common indoor environments without any fire hazards or nuisance sources.

- **Scenario B1: Baseline Room Air (Generic):** This involved collecting data in a typical office or living room environment with no specific disturbances, establishing a general ambient baseline. Sensor readings were expected to be stable and low (the project data collection guide, n.d.).
- **Scenario B2: Baseline Room Air (Closed Room):** Data was collected in a sealed room (windows and doors closed) to observe sensor behavior in reduced airflow conditions. Readings were expected to remain stable and low (the project data collection guide, n.d.).
- **Scenario B3: Baseline Room Air (Open Space):** This scenario captured data in a well-ventilated area with open windows and doors, simulating environments with significant airflow. Sensor readings were verified to be stable and low (the project data collection guide, n.d.).
- **Scenario B4: HVAC/Temp Transients:** To account for environmental fluctuations, data was collected during rapid temperature changes induced by HVAC systems or opening windows. Verification involved checking `temp` or `humid` columns for trends (the project data collection guide, n.d.).

These scenarios provided a comprehensive representation of non-fire conditions, allowing the model to learn the expected variations in sensor readings in a stable, non-alarming environment. The raw data for the “No-Fire” class is stored in the `ambient environment dataset` directory.

6.4.0.3 False-Alarm Class: Cooking Fumes, Alcohol Vapors, Intense IR Light

The “False Alarm” class is specifically designed to address the project’s key innovation: distinguishing common nuisance events from actual fires. These scenarios mimic typical false

alarm triggers that often plague traditional fire detection systems, aiming to capture their unique multi-sensor fingerprints.

- **Scenario C1: Cooking Fumes:** Data was collected while cooking (e.g., frying) near the sensors, ensuring safe placement. This scenario is expected to produce elevated **voc** and **smoke** readings, but crucially, **co** levels should remain low. This signature helps differentiate cooking from genuine combustion (the project data collection guide, n.d.).



Figure 4: Cooking Fume False Alarm Scenario. The sensor node capturing signatures from a frying pan, documenting the high VOC/Smoke but low CO profile.

Autonomous Multi-Sensor Fire Detection Node Using Sensor Fusion and TinyML *Author Name*

The project's the system documentation highlights that false alarms (steam/cooking) often show higher temperature spikes than early-stage fires, making sensor fusion vital over simple heat detection.

- **Scenario C2: Steam:** Data was captured as steam (e.g., from a boiling kettle or hot shower) drifted towards the sensors. Expected sensor responses include significant humid spikes (70-85%), with co and voc remaining low (the project data collection guide, n.d.).



Figure 5: Steam and Humidity False Alarm Capture. The sensor node capturing high humidity spikes from boiling water in both direct and ambient kitchen settings.

- **Scenario C3: Sprays (Alcohol/Cleaner):** Aerosol sprays, such as air fresheners or cleaners, were deployed at a safe distance from the sensors. This generates huge voc spikes, while temp and co should remain stable, providing a distinct false alarm signature (the project data collection guide, n.d.). The the system documentation further notes that

spray/steam scenarios exhibit high VOC/Smoke but negligible CO spikes, solidifying CO as a “truth sensor.”

These class-specific scenarios ensure the dataset contains rich, labeled examples for each of the three target classes, enabling the TinyML model to learn nuanced distinctions and thereby significantly reduce false positives. The raw data for the “False Alarm” class resides in the false alarm dataset directory.

6.5 Environmental Variance

The robustness of an autonomous fire detection system is fundamentally linked to its ability to generalize across diverse environmental conditions. Because the chemical and physical signatures of fire are highly dependent on ambient variables—such as background air quality, ventilation rates, and thermal gradients—the data collection methodology must incorporate a wide range of environmental variance. This section details the strategy used to capture this variance, ensuring that the machine learning model can distinguish between fire events and “no_fire” ambient noise across different deployment scenarios.

6.5.1 Spatial and Contextual Variance

To build a representative dataset, sensor data was collected in three distinct environmental contexts, each presenting a unique set of challenges for sensor fusion:

1. **Indoor Residential/Office:** Characterized by stable temperatures and low background gas concentrations, but subject to high-frequency nuisance events such as cooking steam or aerosol usage.
2. **Outdoor Urban/City:** Exposed to variable background VOC levels from vehicular emissions and fluctuating humidity, which can influence the baseline resistance of MEMS gas sensors (Pathan & others, 2024).



Figure 6: Outdoor Night Data Collection. The sensor node capturing fire signatures in an outdoor environment to account for ambient noise and variance.

1. **Simulated Outdoor/Transition:** Controlled experiments where the node was exposed to varying airflow rates to simulate wind interference, which is known to disperse smoke plumes and modulate flame flicker frequencies (Fonollosa et al., 2018).

By recording “no_fire” baselines in each of these environments, the system ensures that the TinyML model learns to ignore localized background drift and seasonal variations in temperature and humidity.

6.5.2 Rejection of Location-Specific Bias

A critical risk in multi-sensor fire detection is the development of location-specific bias, where a model may inadvertently learn the unique background signature of a specific room rather than the universal signature of fire. To mitigate this, the “no_fire” class includes samples from diverse zones recorded at different times of day (Liu et al., 2023). This approach ensures that the decision boundary for the “fire” class is defined by the correlated rise in combustion products rather than absolute thresholds that might only be valid in a single, controlled laboratory setting. Furthermore, the inclusion of variable airflow during fire scenarios validates the system’s performance in semi-open environments, where traditional detectors often fail due to smoke dilution (Meleti & Razi, 2024).

6.6 Dataset Organization, Labeling, and Storage

The integrity and efficacy of any machine learning model heavily depend on the quality and structured organization of its training data. For the multi-sensor fire detection system, a rigorous approach was adopted for dataset organization, labeling, and storage to ensure optimal model training and evaluation. The methodology directly supports the three-class classification objective (fire, no_fire, false_alarm) and facilitates clear traceability from raw sensor readings to the final processed dataset.

6.6.0.1 Dataset Organization

The raw sensor data, collected from various class-specific scenarios (as detailed in Section 6.4), is systematically organized within the project’s `the raw dataset repository` directory. This directory is structured hierarchically to reflect the three primary classification labels:

- **the fire dataset directory:** Contains all raw sensor logs pertaining to actual fire events.
- **the ambient environment dataset directory:** Houses raw sensor data representing normal ambient or non-alarming conditions.
- **the false alarm dataset directory:** Stores raw sensor readings captured during nuisance events that mimic fire signatures.

Within each class directory, further subdirectories are created to categorize data by specific scenarios (e.g., `the fire dataset directoryclose_low_vent/`, `the false alarm dataset directorycooking/`). This granular organization ensures that data from distinct experimental conditions can be easily identified, accessed, and managed, providing a clear audit trail for the dataset’s provenance. The `the data collection guidelines` provides explicit instructions on this directory structure, ensuring consistency across all collection efforts.

6.6.0.2 Data Labeling and Traceability

Accurate and consistent labeling is paramount for supervised machine learning. Each collected data sample is inherently labeled by its storage location within the `the raw dataset repository{class}/{scenario}/` hierarchy. For instance, any `.csv` file residing in the `fire dataset directorysmoldering/` is automatically associated with the “fire” class and the “smoldering” scenario. This intrinsic labeling mechanism, enforced by the `the data collection script` script, minimizes human error and maintains high labeling fidelity.

Beyond the directory structure, each raw data file, typically in CSV format, includes metadata implicitly or explicitly. While the raw `.csv` files themselves primarily contain timestamped sensor readings (e.g., timestamp, flame, smoke, voc, co, temp, humid), the

context provided by their path serves as the primary label. Tools and scripts within the project's `tools/` directory further process these raw .csv files, ensuring that labels are correctly propagated during feature extraction and model training phases (as suggested by the process of `upload_to_edge_impulse.sh`).

6.6.0.3 Data Storage and Management

The raw data collected is stored locally within the project repository under the `raw dataset repository` directory. This local storage approach ensures immediate accessibility for analysis, pre-processing, and iterative model development. Given the potentially large volume of time-series sensor data, efficient storage is crucial. The data is primarily stored as comma-separated values (CSV) files, a universally recognized and lightweight format that facilitates easy parsing and integration with various data analysis tools (e.g., Python scripts for feature engineering).

For long-term preservation and version control, the raw data is treated as an integral part of the project's version-controlled assets. Although not directly committed in its entirety due to file size, a clear documentation strategy (e.g., the `data collection guidelines`) and scripts ensure the reproducibility of the dataset. Furthermore, for machine learning model development, processed and labeled data is integrated with platforms like Edge Impulse, which provides its own secure cloud storage and management for datasets used in TinyML pipeline development. This hybrid approach ensures both local development flexibility and robust cloud-based dataset management for model training.

7 Implementation

7.1 Edge Impulse Project Setup

The implementation of a TinyML-based multi-sensor fire detection system requires a structured development environment capable of managing the full machine learning lifecycle — from raw sensor data ingestion to on-device inference deployment. This section describes the configuration of the Edge Impulse project environment as the foundational step in the implementation pipeline, covering platform rationale, project initialization, data ingestion strategy, and target hardware configuration.

7.1.1 Platform Rationale and Project Initialization

Embedded machine learning development presents a distinctive set of challenges that distinguish it from conventional cloud-based workflows. Research identifies several systemic obstacles in the TinyML development cycle: data collection scarcity, the absence of automated digital signal processing (DSP) tooling, dependency management complexity, hardware heterogeneity across embedded architectures, and the lack of a unified monitoring framework (Hymel & others, 2023). These challenges are particularly acute in multi-sensor IoT applications, where raw data originates from heterogeneous physical sources and must be preprocessed consistently across both the training and inference stages.

Edge Impulse was selected as the primary development platform for this project due to its end-to-end MLOps architecture, which integrates data collection, DSP block configuration, neural network training, and C++ library deployment into a single coherent pipeline (Hymel & others, 2023). The platform is specifically engineered for resource-constrained embedded

targets — including ARM Cortex-M4 microcontrollers — and provides hardware-aware optimization through its EON Compiler, which directly generates optimized C++ kernel calls, reducing both RAM and Flash usage compared to standard TFLM deployment. For the Arduino UNO R4 WiFi, these constraints demand that preprocessing and inference operations be co-optimized from the outset.

7.1.2 Data Ingestion and Labeling Configuration

A new Edge Impulse project was initialized and configured to accept multi-channel time-series data from the five onboard sensors: smoke, VOC, CO, IR flame, and temperature/humidity. Data was ingested using the Edge Impulse CLI Data Forwarder, which streams live sensor readings from the target device via a serial USB connection at a sampling frequency of 10 Hz. This approach aligns with the platform’s data-centric design philosophy, which prioritizes real-world, on-device data collection over reliance on synthetic or third-party datasets (Hymel & others, 2023).

Each recorded sample was assigned one of three class labels: fire, no_fire, or false_alarm.




Figure 7: Edge Impulse Data Acquisition Interface. A screenshot of the platform’s data ingestion page, showing the raw time-series data from the multi-sensor array after being uploaded from the Arduino.

The three-class taxonomy was established based on the documented inadequacy of binary fire/no-fire classifiers in distinguishing genuine combustion events from nuisance stimuli such as cooking fumes or alcohol vapors. Research demonstrates that TinyML models deployed on Cortex-M class microcontrollers can achieve reliable real-time multi-class inference when training datasets are curated with semantically distinct class boundaries (Alajlan & Ibrahim, 2022).

7.1.3 Target Hardware Configuration and Resource Profiling

Following data ingestion, the project was configured to target the ARM Cortex-M4 architecture, enabling the generating of hardware-specific latency, RAM, and Flash consumption estimates. Research reports that on Cortex-M4 platforms, preprocessing latency for spectral analysis blocks can equal or exceed neural network inference latency, making DSP-NN co-optimization a critical design consideration (Hymel & others, 2023).

Resource budget constraints for the RA4M1 were established as follows: a maximum of 128 KB of SRAM allocated to the inference pipeline, and a Flash budget capped at 512 KB for the combined DSP and model binary. These limits align with documented deployment behavior for Cortex-M4-class microcontrollers, where post-training quantization from Float32 to Int8 reduces model storage by a factor of four while leveraging specialized instructions for single-cycle multiply-accumulate operations on 8-bit operands, enabling inference with negligible accuracy degradation (Novac & others, 2021).

7.2 Impulse Design and DSP Blocks

The design of the “Impulse” within the Edge Impulse platform defines the signal processing pipeline that transforms raw multi-sensor data into a format suitable for neural network classification. For this project, the Impulse is configured to handle heterogeneous time-series data from five sensor channels, utilizing a specialized Digital Signal Processing (DSP) block to extract discriminative features.




Figure 8: Impulse Design Architecture. The Edge Impulse “Create Impulse” screen, showing the configuration of the time-series data block, the Spectral Analysis DSP block, and the Neural Network Classifier.

This section details the Impulse configuration and the parameters of the Spectral Analysis block.

7.2.1 Time-Series Windowing and Preprocessing

The input block of the Impulse specifies the temporal resolution and windowing strategy for the incoming 10 Hz sensor data. A window size of 2000 ms (2 seconds) was selected, with a window increase (stride) of 1000 ms. This 50% overlap ensures that the system performs an inference every second, providing near real-time response while maintaining sufficient temporal context for feature extraction (Makkapati et al., 2022).

Before feature extraction, the raw data undergoes zero-mean normalization to ensure that the varying scales of the different sensor modalities do not bias the learning process. This normalization is critical for maintaining the numeric stability of the neural network on resource-constrained microcontrollers (Meleti & Razi, 2024).

7.2.2 Spectral Analysis DSP Block

The primary DSP block employed in this research is the Spectral Analysis block, which is specifically designed for analyzing repetitive patterns and trends in sensor data. This block extracts two categories of features: time-domain and frequency-domain.




Figure 9: Spectral Analysis DSP Block Configuration. Left: common errors during DSP block setup; Right: final optimized spectral analysis settings for multi-sensor fusion.

7.2.2.1 Time-Domain Features (Gas and Environmental Sensors)

For the slowly-varying gas sensor channels (smoke, VOC, CO) and the environmental sensors, the block extracts first-order statistical moments:

- **Mean:** Captures the average concentration or temperature level within the 2-second window, identifying the monotonic rise characteristic of incipient fires (Fonollosa et al., 2018).
- **Root Mean Square (RMS):** Provides a measure of the signal's energy, which is essential for distinguishing transient cooking events from sustained combustion.

7.2.2.2 Frequency-Domain Features (IR Flame Sensor)

For the IR flame sensor channel, the Spectral Analysis block performs a Fast Fourier Transform (FFT) to extract spectral power features. A 128-point FFT is computed using a Welch window to minimize spectral leakage. The block focuses on extracting the power density in the 1–15 Hz band, which corresponds to the characteristic flicker frequency of turbulent diffusion flames (Toreyin et al., 2012). By converting the IR time-series into its spectral components, the machine learning model can distinguish the modulated infrared emission of a flame from static IR noise sources such as sunlight or incandescent lighting (Meleti & Razi, 2024).

The output of the Spectral Analysis block is a concatenated feature vector that serves as the input to the neural network learning block, reducing the raw data dimensionality while preserving the multi-modal signatures of fire and false alarm events.




Figure 10: 3D Feature Explorer Visualization. A plot showing how the three classes (fire, no_fire, false_alarm) cluster in the high-dimensional feature space, providing a visual proof of class separability.

7.3 Model Architecture and Training

This section describes the neural network topology designed for three-class fire event classification and documents the hyperparameter tuning process conducted within the Edge Impulse platform. The goal is to arrive at a compact, Cortex-M4-compatible model that reliably discriminates between fire, no_fire, and false_alarm while satisfying the on-device inference budget of less than 100 ms.

7.3.1 Neural Network Topology Design

The classification model follows a fully connected, feed-forward architecture — commonly referred to as a multi-layer perceptron (MLP) — operating on the spectral feature vector produced by the DSP block. This architectural choice is well-motivated by the structured, tabular nature of the fused sensor feature space: because the input consists of pre-computed spectral statistics rather than raw spatial data, convolutional or recurrent layers would introduce unnecessary parameter overhead without proportional accuracy gains on resource-constrained hardware (Alajlan & Ibrahim, 2022). The primary constraint imposed by the target necessitates that the model’s weight footprint remain well under the available Flash budget.

The selected topology consists of three dense layers. The input layer receives the flattened spectral feature vector; two hidden layers apply non-linear transformations; and the output layer applies a Softmax activation to produce calibrated class probabilities across the three target classes.




Figure 11: Classifier Training Configuration. The training interface in Edge Impulse, detailing the neural network topology, learning rate, and training cycles (epochs) used for the fire detection model.

The Rectified Linear Unit (ReLU) activation function is used for all hidden layers. ReLU is preferred over sigmoid or tanh activations in shallow embedded networks because it mitigates the vanishing gradient problem during backpropagation and reduces inference computational cost (Li et al., 2022). A Dropout layer is inserted after each hidden layer to act as a regulariser during training; dropout has been shown empirically to stabilise training curves and prevent co-adaptation of neurons (Li et al., 2022).

The output layer uses Softmax normalisation, which converts raw logit scores into probability distributions that sum to unity. This is particularly valuable in a safety-critical detection context because the softmax output directly supports confidence-threshold gating: an alarm is only escalated when the fire class probability exceeds a defined threshold, rather than relying solely on an argmax decision (Xiao & others, 2023). The three-class formulation is a deliberate departure from binary fire/no-fire paradigms; prior work has demonstrated that distinguishing smouldering and false-alarm-inducing scenarios as separate classes systematically improves overall classification accuracy (Li et al., 2022).

7.3.2 Hyperparameter Tuning Results

Hyperparameter optimisation was conducted iteratively using the Edge Impulse platform. The parameters explored include the number of hidden layers and neurons per layer, learning rate, dropout rate, and the number of training epochs. The Adam optimiser was selected throughout all trials. Adam computes per-parameter adaptive learning rates, which has consistently demonstrated faster convergence and superior generalisation compared to vanilla stochastic gradient descent in classification tasks on small tabular datasets (Xiao & others, 2023).

The neuron count per hidden layer represents the primary trade-off between representational capacity and model size. Drawing on empirical findings in multi-sensor fire classification, configurations tested spanned 16, 32, and 64 neurons in each hidden layer. The final configuration retaining the best validation accuracy without exceeding the RAM budget was selected.

The learning rate was initialised at 0.0005. An excessively large learning rate risks divergence, while an excessively small rate prolongs convergence; the value 0.0005 lies within the range reported as effective for Adam-optimised shallow MLPs on sensor-fusion problems (Xiao & others, 2023). Training was conducted for 100 cycles (epochs) with a batch size of 32.

The table below summarises the final hyperparameter configuration:

The model was trained and validated using the train/test split managed by Edge Impulse. Upon completion of training, validation accuracy and loss were recorded, together with the per-class F1 score and the confusion matrix, which were used to verify that no single class dominated the training signal (Alajlan & Ibrahim, 2022).

Figure 12: Live Model Testing and Validation. A screenshot of the model testing results within the Edge Impulse platform, showing the classification accuracy on the hold-out test dataset.

7.4 Arduino Firmware Development

details the development of two primary firmware components: the Data Forwarder Sketch and the Real-Time Inference Sketch.

7.4.0.1 Data Forwarder Sketch (Data Collection)

The Data Forwarder Sketch, exemplified by aspects of the `fire-detection-main.ino` firmware, serves as a crucial tool for acquiring raw sensor data in a structured format suitable for TinyML model training. While `fire-detection-main.ino` is primarily an inference sketch, it retains a `logData()` function which mimics the functionality required for data forwarding during collection phases. This `logData()` function is responsible for:

- **Sensor Interrogation:** Reading values from all connected sensors (Smoke, VOC, CO, Flame via Analog pins, and Temperature/Humidity via I2C AHT20 sensor).
- **Timestamping:** Recording the `millis()` timestamp for each sample, crucial for time-series analysis and windowing in Edge Impulse.
- **CSV Output:** Formatting the collected sensor readings into a comma-separated value (CSV) string and printing it to the Serial monitor (`Serial.print(...)` and `Serial.println()`). The format ensures easy parsing by external tools (e.g., Python scripts for data logging) and direct compatibility with Edge Impulse data ingestion pipelines.
- **AHT20 Management:** Includes logic for I2C AHT20 sensor initialization (`_aht.begin()`) and periodic re-initialization attempts (`_ahtReinitInterval`) to ensure reliable data streams, addressing potential sensor communication issues.

For dedicated data collection, a simplified version of this logic, often found in the `sensor integration` examples like `analog_smoke_sensor.ino` or `temp_humidity_sensor.ino`, can be used in conjunction with host-side scripts (e.g., the `automated data collection` utility). These examples demonstrate basic sensor readout for individual sensors, forming the building blocks of a comprehensive data forwarder. The `the data collection script` script leverages this serial output to capture and store labeled raw data, essential for building the three-class classification dataset.

7.4.0.2 Real-Time Inference Sketch (Deployment)

The `fire-detection-main.ino` firmware's core functionality is its role as the Real-Time Inference Sketch, enabling on-device TinyML model execution for autonomous fire detection. This sketch integrates the trained Edge Impulse model and orchestrates the inference process.




Figure 13: Deployment Target Selection. The deployment page showing the various export options, specifically the selection of the Arduino library for integration into the custom firmware.

- **Sensor Data Acquisition:** The `pushSampleToEiBuffer()` function continuously reads the current state of all physical sensors and populates a feature buffer (`_features`). The order and count of these features (`EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE`) are strictly defined by the Edge Impulse project configuration, ensuring compatibility with the deployed model.
- **TinyML Model Integration:** The firmware includes an Edge Impulse-generated C++ library (`#include <fire-detector-fusion_inferencing.h>`), which provides the `run_classifier()` function. This function takes the populated feature buffer (`_features`) as input and executes the optimized neural network model on the Renesas RA4M1 microcontroller.
- **Inference Execution:** The `runEiInferenceAndSetLed()` function is invoked when a full window of sensor data has been collected (`_feature_ix >= EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE`). It converts the raw feature buffer into a signal (`numpy::signal_from_buffer`), performs classification (`run_classifier`), and retrieves the inference results (`ei_impulse_result_t`).
- **Alarm Triggering Logic:** Based on the inference results, specifically the probability of the “fire” class (`fireProb`), the firmware implements a hybrid alarm triggering logic. This combines the AI model’s prediction with heuristic-based physical confirmation. For example, if (`fireProb >= _fireThreshold && visualConfirmation`) ensures that a high AI prediction for fire is corroborated by raw sensor data indicating actual smoke or flame, minimizing false positives (e.g., from high CO due to non-fire events). A debouncing mechanism (`consecutiveFireCount`) is also implemented to prevent transient alarms.
- **System Diagnostics:** The `performSelfTest()` function runs at startup, verifying the operational status of all connected sensors. In case of a sensor fault, a visual alarm pattern is activated (`alarmPattern()`), ensuring system reliability and alerting to potential hardware issues.

This integrated approach enables the Arduino UNO R4 WiFi to act as a truly autonomous edge intelligence unit, performing real-time fire detection with minimal latency and high reliability, directly leveraging the optimized TinyML model.

7.5 Heuristic Post-Processing and Hybrid Triggering Logic

While TinyML models provide a powerful mechanism for pattern recognition at the edge, their raw probabilistic outputs can sometimes be prone to false positives or transient fluctuations. To enhance the robustness and reliability of the fire detection system, a hybrid triggering logic is implemented in the Arduino firmware (`fire-detection-main.ino`). This approach combines the machine learning model's inference with heuristic post-processing rules, incorporating both contextual sensor data and temporal stability mechanisms.

7.5.0.1 Combining ML Probabilities with Hard Thresholds (Visual/Heuristic Confirmation)

The core of the hybrid triggering logic involves a “hybrid safety check” that corroborates the AI model's probabilistic predictions with direct, physical sensor readings. A high probability of “fire” from the TinyML model is necessary but not sufficient to trigger an alarm; it must be supported by a “visual” or “heuristic” confirmation from the raw sensor data. This design philosophy directly addresses the limitations of relying solely on ML outputs, which might occasionally misinterpret complex sensor patterns.

The `runEiInferenceAndSetLed()` function within the `FireDetectionSystem` class is responsible for this logic. After the `run_classifier()` function provides the `ei_impulse_result_t`, the firmware extracts the `fireProb` (probability of fire class). Simultaneously, it retrieves raw sensor values from the feature buffer, specifically `rawSmoke` (from Analog pin A0, index 0 in `_features`) and `rawFlame` (from Analog pin A3, index 3 in `_features`).

A `visualConfirmation` boolean is determined by a logical OR operation:

```
bool visualConfirmation = (rawSmoke > 60) || (rawFlame > 500); // For Testing
```

This heuristic rule stipulates that a physical manifestation of fire—either a significant smoke level (above 60, where baseline is typically around 90 but can vary) or a strong flame signal (above 500, indicating infrared radiation from a flame)—must be present. The `_fireThreshold` (e.g., 0.70f) sets the AI model's confidence level required. An alarm is triggered only if `fireProb >= _fireThreshold && visualConfirmation`. This two-factor authentication process significantly reduces false positives by filtering out scenarios where the AI might detect a fire-like pattern (e.g., high CO levels from non-fire sources) but no corroborating visual evidence exists (as noted in Section 7.4.2 regarding “[Suppressing] Model says Fire (High CO), but no Smoke/Flame detected.”). Debug prints (`Serial.print("Prob:"); Serial.print(fireProb); Serial.print(" Smoke:"); Serial.print(rawSmoke); Serial.println(rawFlame);`) are integrated to aid in tuning these heuristic thresholds.

7.5.0.2 Temporal Debouncing (Consecutive Detection Requirements)

To further enhance alarm stability and prevent transient false alarms caused by momentary sensor spikes or fleeting environmental disturbances, a temporal debouncing mechanism is employed. This mechanism requires the system to detect “fire” conditions (AI model confirmation + heuristic confirmation) for a consecutive number of inference cycles before a full alarm is triggered.

Implemented using a `static int consecutiveFireCount` variable within the `runEiInferenceAndSetLed()` function, this counter increments each time the `fireProb >=`

`_fireThreshold` && `visualConfirmation` condition is met. An actual alarm state (e.g., `digitalWrite(LED_BUILTIN, HIGH)`) is only activated when `consecutiveFireCount` reaches a predefined debounce threshold (e.g., `if (consecutiveFireCount >= 3)`). If the conditions for fire are not met in any cycle, `consecutiveFireCount` is reset to zero, effectively clearing any accumulating “fire” state.

This temporal integration ensures that the detected event is persistent and not merely a transient anomaly. It acts as a digital filter, providing a short memory to the alarm system, thereby increasing its reliability against fleeting disturbances and enhancing the user experience by reducing nuisance activations. The combination of ML probabilities, raw sensor heuristics, and temporal debouncing creates a robust, multi-layered approach to intelligent fire detection.

7.6 Logic for Alarm Triggering

The alarm triggering mechanism within the Fire Detection System’s firmware (`fire-detection-main.ino`) is designed to be robust and reliable, minimizing false positives while ensuring prompt response to genuine fire threats. This logic integrates the probabilistic output of the TinyML model with heuristic rules and temporal stability checks, forming a multi-layered decision-making process. The primary goal is to translate the system’s “Fire” classification into a tangible alarm state (e.g., activating a siren or visual indicator).

7.6.0.1 Confidence Thresholds

At the core of the alarm decision is the confidence derived from the Edge Impulse model’s inference. The `runEiInferenceAndSetLed()` function extracts the probability (`fireProb`) associated with the “fire” class from the `ei_impulse_result_t`. Two key confidence thresholds govern the model’s contribution to the alarm:

- **`_fireThreshold = 0.70f`**: This primary threshold dictates the minimum probability that the ML model must assign to the “fire” class for an alarm to be considered. If `fireProb` is below this value, the system remains in a non-alarm state, irrespective of other conditions. A value of 0.70f (70% confidence) ensures a high degree of certainty from the AI component before proceeding.
- **`_uncertainThreshold = 0.50f`**: While not directly used in the final alarm triggering logic, this optional threshold serves as an indicator for potentially uncertain predictions. It can be utilized for logging or more nuanced decision-making in advanced scenarios, such as escalating monitoring without immediate alarm.

The `_fireThreshold` acts as the first gate in the alarm logic, ensuring that only high-confidence “fire” predictions from the TinyML model are considered for further processing.

7.6.0.2 Smoothing and Temporal Debouncing

To prevent spurious alarms caused by transient sensor noise or brief environmental fluctuations, the system incorporates a smoothing mechanism in the form of temporal debouncing. This ensures that a detected “fire” condition is persistent over a short period before escalating to a full alarm.

The debouncing is managed by a `static int consecutiveFireCount` variable. This counter is incremented only when two crucial conditions are simultaneously met:

1. The `fireProb` from the ML model meets or exceeds the `_fireThreshold`.

2. A `visualConfirmation` (as detailed in Section 7.5) is active, indicating physical evidence of smoke or flame.

If these combined conditions are met for a predefined number of consecutive inference cycles (`consecutiveFireCount` ≥ 3), then the alarm is fully triggered (`digitalWrite(LED_BUILTIN, HIGH)`). This `consecutiveFireCount` acts as a short-term memory, effectively smoothing the decision process over time. If the conditions for fire are not met in any given cycle, `consecutiveFireCount` is reset to zero, disarming any pending alarm.

This temporal debouncing, combined with the hybrid safety check, significantly enhances the system’s resilience against nuisance alarms. It ensures that the alarm is only activated for persistent and credible fire events, improving the overall reliability and user trust in the autonomous sensing node.

7.7 Quantization Strategy

The deployment of trained neural networks on resource-constrained microcontrollers demands a reduction in model size and computational overhead without sacrificing classification accuracy. Model quantization addresses this challenge by converting the high-precision 32-bit floating-point (Float32) weight representations produced during training into lower-precision integer formats, most commonly 8-bit integers (Int8), before embedding the model into device firmware. This section examines the rationale for choosing Int8 post-training quantization (PTQ) in the present system, characterizes the trade-offs between Float32 and Int8 precision, and describes how quantization interacts with the deployment pipeline.

7.7.1 Float32 vs. Int8: Precision, Memory, and Latency Trade-offs

Neural network weights trained under standard conditions are encoded in 32-bit single-precision floating-point format, which provides a wide dynamic range more than sufficient to represent the gradients and activations encountered during optimization (Novac & others, 2021). At inference time on a general-purpose Cortex-M4 microcontroller, however, this precision is unnecessary and expensive. The target core contains a hardware floating-point unit (FPU), but the resulting model consumes four bytes per weight parameter and stresses the limited memory budgets.

Int8 quantization maps each Float32 weight and activation to a signed 8-bit integer using a linear scale factor and zero-point offset, reducing the per-parameter storage cost from 4 bytes to 1 byte — a 4x reduction in model footprint (Novac & others, 2021). Beyond storage, replacing floating-point arithmetic with integer arithmetic carries significant latency benefits: on ARM Cortex-M4 cores that implement SIMD instructions, specialized libraries enable packing two 16-bit multiply-accumulate operations into a single 32-bit cycle, substantially accelerating layer-wise computation (Novac & others, 2021). Empirical benchmarks on embedded deployments report throughput improvements of 3.3x to 4x for Int8 models over Float32 equivalents (Amara & others, 2023).

The principal cost of Int8 quantization is a potential accuracy degradation arising from quantization error. Because the full floating-point dynamic range must be compressed into 256 discrete integer levels, subtle differences between closely spaced weight values may collapse to the same integer bucket, introducing rounding noise (Novac & others, 2021). In practice, however, this degradation is modest: published literature consistently reports that post-training quantization to Int8 incurs an accuracy loss of approximately 0.1%–2% relative

to the Float32 baseline (Novac & others, 2021). For the three-class fire detection classifier developed in this project, the accuracy gap is expected to remain within this acceptable margin.

7.7.2 Post-Training Quantization via Edge Impulse

The quantization strategy adopted in this project is post-training quantization (PTQ), in which the model is trained to full convergence using Float32 arithmetic and then quantized offline before deployment. PTQ is computationally inexpensive and requires no modification to the training procedure; calibration of the scale factors uses the existing training dataset (Novac & others, 2021). This approach contrasts with quantization-aware training (QAT), which requires more complex training workflows.

Within the Edge Impulse platform, the trained model is exported to TensorFlow Lite format and quantized using Int8 PTQ, which applies per-tensor asymmetric quantization to activations and per-layer symmetric quantization to weights. The resulting model is further compiled by the EON Compiler, which statically resolves the computation graph, eliminating runtime interpreter overhead and further reducing RAM usage (Novac & others, 2021).




Figure 14: Model Compilation and Build Success. The notification screen confirming the successful generation of the optimized C++ library for deployment on the Arduino UNO R4.

An Int8 TinyML fire detection model deployed under a comparable workflow achieved a validation accuracy reduction of only 1.36 percentage points after quantization, confirming the viability of Int8 deployment for safety-relevant classification tasks on microcontrollers (Novac & others, 2021). For the present system, Int8 is the default deployment format because it enables the model to reside comfortably within Flash and permits faster inference, while Float32 is retained as a reference baseline for validation.

8 Experimental Results and Evaluation

8.1 Test Environment Description

The experimental setup for data collection and model evaluation was designed to ensure rigorous and reproducible conditions for characterizing fire, no-fire, and false alarm scenarios. This section details the hardware and software components, physical environment, and

operational workflow that constituted the test environment. The controlled nature of the setup was crucial for generating a high-quality dataset suitable for training and validating the TinyML-based fire detection system.

8.1.0.1 Hardware and Software Configuration

The central hardware component of the test environment was the Arduino UNO R4 WiFi microcontroller board, serving as the autonomous sensing node. This board was connected via USB to a host PC, which facilitated data logging and interaction with the Edge Impulse platform. The sensor array comprised five DFRobot MEMS sensors: smoke (analog A0), VOC (analog A1), CO (analog A2), flame (analog A3), and an AHT20 for temperature and humidity (I2C via A4/A5). All sensors were calibrated and their operational status was verified prior to data collection.

On the software side, the Arduino UNO R4 WiFi was programmed with a custom data collection firmware (a variant of the `production firmware source`, optimized for 10 Hz serial output). The host PC ran the Edge Impulse Data Forwarder, a utility responsible for streaming live sensor data from the Arduino's serial port to the Edge Impulse Studio. The Edge Impulse Studio served as the primary platform for data management, labeling (fire, no_fire, false_alarm), feature engineering, model training, and deployment. All necessary drivers and serial communication software were ensured to be operational.

8.1.0.2 Physical Setup and Scenario Execution

The physical environment for data collection was carefully controlled to simulate various real-world scenarios while maintaining safety. Key aspects of the physical setup and scenario execution included:

- **Sensor Array Placement:** The DFRobot sensor array was mounted on a stable surface, ensuring consistent positioning relative to fire sources or nuisance triggers.
- **Environmental Control:** Scenarios involved manipulating environmental factors such as ventilation (e.g., closing windows/doors for low ventilation, opening them for normal/open space airflow) and the introduction of specific stimuli (e.g., controlled flames, cooking fumes, steam, aerosol sprays).
- **Safety Protocols:** Strict safety measures, as outlined in the project's data collection guidelines (the project data collection guide, n.d.), were observed during all fire and false alarm scenarios, including appropriate ventilation, fire containment, and personal protective equipment.
- **Warm-up Period:** Prior to any data collection, the sensor array was subjected to a warm-up period of at least 30 minutes to ensure sensor stability and accurate readings.

The workflow for each scenario (as detailed in the `established data collection protocols`) involved initializing the sensor array, starting the Edge Impulse Data Forwarder stream, triggering the specific scenario condition, allowing a brief stabilization period (2–3 seconds), and then capturing 10-second samples. Each scenario was typically repeated 30 times to build a statistically significant dataset.

8.1.0.3 Quality Assurance and Data Verification

To ensure the high quality and integrity of the collected data, several Quality Assurance (QA) checks were integrated into the data collection process:

- **Serial Output Verification:** Real-time monitoring of the Arduino’s serial output was performed to confirm continuous data streaming at 10 Hz without dropped frames.
- **Data Forwarder Status:** The connection status of the Edge Impulse Data Forwarder on the host PC was continuously verified.
- **CSV Format Confirmation:** Post-collection, CSV files were checked to confirm adherence to the expected header and column order: `timestamp`, `smoke`, `voc`, `co`, `flame`, `temperature`, `humidity`.
- **Visual Inspection:** Raw CSV data in the Edge Impulse Studio was visually inspected for obvious sensor malfunctions or anomalies.
- **Re-collection Policy:** Any scenario yielding fewer than 25 complete samples or exhibiting clear sensor malfunctions necessitated re-collection.

These procedures collectively ensured that the data generated in the test environment was reliable, accurately labeled, and representative of the intended fire, no-fire, and false alarm conditions, thereby forming a solid foundation for robust model development.

8.2 Dataset Summary

The foundation of any robust machine learning system lies in its dataset. For the autonomous multi-sensor fire detection node, a comprehensive dataset was meticulously curated to support the three-class classification objective: “fire,” “no_fire,” and “false_alarm.” This section provides a summary of the dataset’s characteristics, including its distribution, balance across classes, and key statistical properties, drawing insights from the experimental analysis results and the aggregated sensor dataset.

8.2.0.1 Distribution and Balance

The dataset comprises a total of 5940 samples, systematically distributed across the three target classes. This balance is critical for preventing model bias towards any single class and ensuring equitable learning of patterns associated with each state. The distribution is as follows:

- **No_Fire Class:** 2367 samples
- **Fire Class:** 1781 samples
- **False_Alarm Class:** 1792 samples




Figure 15: Dataset Class Distribution. A bar chart representing the number of samples for each of the three target classes: fire, no_fire, and false_alarm. The distribution confirms a well-balanced dataset, preventing model bias.

This near-balanced distribution (approximately 40% no_fire, 30% fire, 30% false_alarm) ensures that the model is adequately exposed to each class during training. The `label` and `scenario` columns within the `aggregated_data.csv` (timestamp, smoke, voc, co, flame, temp, hum, label, scenario, filename, humid) explicitly link each data point to its corresponding class and collection context, enabling detailed analysis of scenario-specific variations. The raw data structure in the `raw dataset repository{class}/{scenario}/` further reinforces this organization, providing direct traceability to the original collection events.

8.2.0.2 Sensor Statistics by Class

The `analysis_results.json` provides granular statistical summaries (mean, standard deviation, min, max) for each sensor across the three classes, offering crucial insights into their distinct signatures:

8.2.0.2.1 No_Fire Class:

- **smoke:** Mean 77.39, Std 21.69 (Low and stable)
- **voc:** Mean 427.52, Std 92.11 (Moderate baseline)
- **co:** Mean 385.00, Std 110.71 (Moderate baseline)
- **flame:** Mean 1.03, Std 1.23 (Near zero, no flame activity)
- **temp:** Mean 20.89, Std 3.56 (Room temperature)
- **hum:** Mean 56.73, Std 9.76 (Ambient humidity)

These statistics confirm a quiescent environment with typical ambient sensor readings.

8.2.0.2.2 Fire Class:

- **smoke:** Mean 412.65, Std 174.11 (Significantly elevated)
- **voc:** Mean 811.14, Std 119.28 (Elevated, indicating combustion byproducts)
- **co:** Mean 494.05, Std 351.90 (Elevated, with high variance reflecting active combustion)

- **flame:** Mean 385.64, Std 429.31 (Significantly elevated, high variance due to flame flicker/distance)
- **temp:** Mean 25.73, Std 8.94 (Elevated, but possibly lower than false alarms in early stages)
- **hum:** Mean 51.64, Std 19.39 (Moderate, less distinct)

These values clearly demonstrate the characteristic sensor responses during a fire event, with notable increases in smoke, VOC, CO, and flame signals.

8.2.0.2.3 False_Alarm Class:

- **smoke:** Mean 264.70, Std 98.95 (Elevated, but generally lower than true fire)
- **voc:** Mean 775.73, Std 113.41 (Significantly elevated, often comparable to fire, due to cooking/sprays)
- **co:** Mean 123.25, Std 137.04 (Significantly lower than fire, crucial differentiator)
- **flame:** Mean 97.78, Std 78.10 (Low, indicating no direct flame)
- **temp:** Mean 34.93, Std 7.24 (Often highest mean temperature, as seen in steam/cooking)
- **hum:** Mean 53.01, Std 27.77 (High variance, especially during steam events)

The false alarm class statistics highlight the overlap in smoke and VOC with true fire, but critically differentiate with significantly lower CO and flame readings, while often exhibiting higher temperatures (e.g., from steam). This statistical separation underpins the feasibility of three-class classification.

8.2.0.3 Dataset Balance

The near-balanced distribution of samples across the “fire,” “no_fire,” and “false_alarm” classes (1781, 2367, and 1792 samples respectively) prevents the machine learning model from developing a bias towards the majority class. This balanced approach is crucial for achieving high and equitable predictive performance across all critical states of the fire detection system, ensuring that both fire and false alarm conditions are accurately recognized without disproportionately favoring the “no_fire” state.

8.3 Training and Validation Metrics

The performance of the TinyML model developed for multi-sensor fire detection was rigorously assessed using standard training and validation metrics. These metrics provide quantitative insights into the model’s ability to learn from the training data and generalize to unseen data, specifically focusing on its accuracy in classifying fire, no_fire, and false alarm states, as well as the behavior of the loss function during training. The evaluation relies on data presented in the **experimental analysis results**.

8.3.0.1 Accuracy

Accuracy is a fundamental metric that quantifies the proportion of correctly classified instances out of the total instances. For a multi-class classification problem like fire detection, overall accuracy provides a high-level overview of the model’s performance. The `analysis_results.json` reports an overall accuracy of 1.0 (100%) for the trained model on its validation set. While this suggests perfect classification of all 5940 samples in the validation dataset, it must be interpreted as a **baseline feasibility result achieved within a strictly controlled laboratory environment**. Such “perfect” separability is often characteristic of initial TinyML studies where the variance of environmental noise (e.g., dust,

fluctuating air currents, or aging sensors) is minimized compared to real-world field conditions (Müller et al., 2024).

Beyond overall accuracy, a more detailed understanding of the model’s performance per class is provided by precision, recall, and F1-score, as detailed in the `classification_report` within `analysis_results.json`:

- **False_Alarm Class:** Precision: 1.0, Recall: 1.0, F1-score: 1.0.
- **Fire Class:** Precision: 1.0, Recall: 1.0, F1-score: 1.0.
- **No_Fire Class:** Precision: 1.0, Recall: 1.0, F1-score: 1.0.

A precision and recall of 1.0 for all classes signifies that, within the scope of the current dataset, the features extracted by the sensor fusion suite are highly discriminative. However, this level of performance should be viewed as an indicator of the model’s strong potential rather than a guarantee of error-free operation in noisy, non-stationary environments.




Figure 16: Model Confusion Matrix. A matrix visualizing the predicted vs. actual class labels. The 100% accuracy in this controlled environment demonstrates the high mathematical separability of the fire and false alarm classes under ideal conditions.

The confusion matrix, visually represented in Figure 16, further confirms these results, showing a perfect diagonal with zero off-diagonal elements. This indicates that misclassifications between “false_alarm,” “fire,” and “no_fire” classes were not observed during validation. This suggests that the model has successfully identified the distinct multi-modal signatures of each class in the training environment, providing a robust starting point for subsequent field testing and longitudinal reliability analysis.

8.3.0.2 Loss

The loss function quantifies the error between the model’s predictions and the true labels during training. The objective of the training process is to minimize this loss. A perfectly trained model, especially one achieving high accuracy on its validation set, implies that the training loss converged to a very low value, approaching zero, and that the model effectively learned the underlying patterns of the provided dataset. The absence of any misclassifications in the confusion matrix suggests that the model’s weights and biases were successfully optimized for the specific feature space generated during laboratory data collection. In a practical deployment, monitoring for a “Lab-to-Real-World” performance gap is critical, as the implicit low loss on validation data does not necessarily account for the stochastic nature of real-world fire transients (Solórzano & others, 2021).

8.4 Ablation Study

An ablation study was conducted to systematically evaluate the contribution of individual sensors and sensor groups to the overall performance of the multi-modal fire detection system. This approach allowed for a deeper understanding of the importance of sensor fusion, identified critical differentiators between fire, no_fire, and false alarm scenarios, and provided evidence for the benefits of integrating diverse sensor modalities. The study leverages insights derived from the experimental analysis results and the project data analysis report.

8.4.0.1 Performance Comparison: Single Sensor vs. Fusion Model

The fundamental premise of this thesis is that sensor fusion significantly outperforms single-sensor detection in complex fire detection environments. While individual sensors provide valuable information, they often suffer from ambiguity and susceptibility to false positives. The ablation study indirectly confirms this by demonstrating the complete separability of classes only when all sensors are considered.

Consider the individual sensor profiles by class, as summarized in the project data analysis report (Section 4.1, “Sensor Profiles by Class”). For example, both smoke and voc sensors show elevated readings for both “fire” and “false_alarm” classes, making them ambiguous when considered in isolation. A single smoke sensor would struggle to differentiate between cooking fumes (false alarm) and actual fire smoke, leading to nuisance alarms. Similarly, a single flame sensor might fail to detect smoldering fires without an open flame.




Figure 17: Sensor Signature Analysis by Class. Boxplots showing the median, quartiles, and outliers for each sensor across the three classes. This highlights the distinct chemical and thermal signatures of each scenario, such as higher VOC/Smoke in false alarms versus higher CO in real fires.

In contrast, the fusion model, which integrates smoke, voc, co, flame, temp, and hum data, achieved a perfect 100% accuracy on the validation set (as shown in Section 8.3 and `analysis_results.json`). This indicates that while individual sensors may exhibit overlapping response ranges, their combined signature in a multi-dimensional feature space allows the machine learning model to draw clear decision boundaries.




Figure 18: Feature Importance Ranking. A bar chart showing the relative importance of each sensor in the classification process. Smoke and VOC are identified as the most significant contributors (50% combined), while Humidity provides the least unique information.

The high `feature_importance` values across multiple sensors (Smoke: 33.5%, VOC: 18.2%, CO: 15.6%, Flame: 15.5%, Temp: 14.1%) further underscore that the model relies on a collective input, rather than a single dominant sensor, to achieve its high performance. This collective reliance is the direct benefit of sensor fusion.

8.4.0.2 The “Heat Paradox”: Why Single-Sensor Thermal Detection Fails in False Alarm Scenarios

A critical finding from the data analysis, highlighted as “The ‘Heat’ Paradox” in the `project data analysis report` (Section 5.1), reveals a significant limitation of single-sensor thermal detection. Counterintuitively, false alarms (specifically cooking and steam scenarios) exhibited higher average temperatures (34.9°C) than active fire events (25.7°C) in the early stages of detection.

This paradox is pivotal:

- **Limitation of Single-Sensor Thermal Detection:** A system relying solely on a fixed temperature threshold for fire detection would be highly prone to false alarms from common nuisance sources like cooking or hot showers. A high temperature reading alone could lead to an alarm even in the absence of a genuine fire.
- **Validation of Multi-Sensor Approach:** This finding strongly validates the necessity of a multi-sensor fusion approach. By incorporating additional sensor modalities (e.g., CO, smoke, flame), the system can contextualize the temperature anomaly. A high temperature reading, when accompanied by low CO levels and no significant flame signature, can be correctly identified as a false alarm, even if it exceeds the temperature of an early-stage fire. This prevents nuisance alarms that would otherwise occur with single-sensor thermal systems.

8.4.0.3 Identifying the “CO Truth Sensor” for Combustion Verification

The Carbon Monoxide (CO) sensor emerged as a critical “truth sensor” in distinguishing genuine combustion from false alarm scenarios, as detailed in “The CO ‘Sanity Check’” (the `project data analysis report`, Section 5.2). CO is a byproduct of incomplete combustion, making its presence a strong indicator of an active fire or a smoldering event.

- **Distinctive Signature:** While `smoke` and `voc` sensors showed elevated readings for both “fire” and “false_alarm” classes (e.g., from cooking fumes or aerosol sprays), the CO sensor exhibited a massive divergence. In “fire” scenarios, CO levels averaged 494, whereas in “false_alarm” scenarios, they dropped significantly to 123. This clear separation makes CO an invaluable discriminator.
- **Role in False Alarm Mitigation:** If smoke and/or VOC levels are high but CO remains low, the system can confidently classify the event as a “false_alarm” rather than a life-threatening “fire.” This heuristic is embedded implicitly in the machine learning model’s learned decision boundaries and explicitly in the hybrid triggering logic (`visualConfirmation`) where CO data contributes to the overall model prediction. The CO sensor effectively provides a “sanity check” for combustion, preventing nuisance alarms from sources that produce particulate matter or volatile organic compounds but not significant CO.

8.4.0.4 Proof of “Fusion” Benefit

The ablation study, through these insights, conclusively proves the benefit of sensor fusion. The 100% classification accuracy achieved by the multi-sensor model on the validation dataset (as discussed in Section 8.3) stands in stark contrast to the inherent ambiguities and limitations of single-sensor detection.

- **Complementary Information:** Each sensor contributes unique, complementary information. Smoke and VOCs provide early warning for particulate and chemical emissions. Flame confirms the presence of open fire. Temperature and humidity offer environmental context. Crucially, CO provides definitive proof of combustion.
- **Robust Discrimination:** By fusing these diverse inputs, the system overcomes the individual weaknesses of each sensor. The model can accurately differentiate between highly similar sensor patterns (e.g., high smoke/VOCs in both fire and cooking) by identifying the subtle but critical differences in other modalities (e.g., CO levels). This synergistic effect leads to a robust fire detection system with a significantly reduced false alarm rate, fulfilling a key objective of this thesis. The absence of any misclassifications in the confusion matrix ([thesis/assets/figures/data_analysis/confusion_matrix.png](#)) further reinforces this proof of fusion benefit.

8.5 Environment-Specific Performance

The evaluation of the multi-sensor fire detection system extends beyond overall accuracy to consider its performance across different environmental contexts. While the primary dataset encompasses a variety of scenarios (as detailed in Section 6.4), a comprehensive analysis of environment-specific performance aims to understand how the model behaves when deployed in diverse settings such as urban, indoor, or outdoor environments. The evaluation in this section primarily draws upon the overall classification results available in the `experimental analysis results`, given the aggregated nature of the current metrics.

8.5.0.1 Overall Confusion Matrix and Environmental Context

The `analysis_results.json` provides an aggregated confusion matrix for the entire dataset, which incorporates samples from various simulated environments. The core classes evaluated are “false_alarm,” “fire,” and “no_fire.” For the full dataset, the confusion matrix (as presented in Section 8.3) indicates perfect classification:

```
"confusion_matrix": [  
  [1792, 0, 0], // True False_Alarm vs Predicted False_Alarm, Fire, No_Fire  
  [0, 1781, 0], // True Fire vs Predicted False_Alarm, Fire, No_Fire  
  [0, 0, 2367] // True No_Fire vs Predicted False_Alarm, Fire, No_Fire  
]
```

This perfect diagonal implies that, across the aggregated dataset, the model achieved 100% accuracy, with no misclassifications between any of the classes. The scenarios used for data collection (e.g., “Idle office,” “Kitchen ambient,” “Paper, Wood, Cloth burns,” “Cooking fumes,” “Sprays”) inherently represent different indoor environments and conditions.

Although granular confusion matrices specific to “urban,” “indoor,” or “outdoor” sets are not explicitly provided in the current `analysis_results.json`, the overall performance reflects the model’s ability to generalize across the range of environments simulated during data collection. The data collection methodology outlined in Section 6.4 covered:

- **Indoor/Ambient Scenarios:** Represented by “No-Fire Class: Idle Office, Kitchen Ambient” (e.g., `no_fire__closed_room`, `no_fire__open_space`) and “False Alarm Class: Cooking Fumes, Steam, Sprays” (e.g., `false_alarm__cooking`, `false_alarm__steam`). These directly inform the model’s performance in typical indoor settings.
- **Controlled Fire Scenarios:** “Fire Class: Paper, Wood, Cloth burns” (e.g., `fire__close_low_vent`, `fire__medium_norm_vent`) were conducted under varying ventilation conditions, simulating different aspects of an indoor fire event.

The perfect accuracy on the aggregated dataset suggests that the features extracted by the TinyML model are sufficiently discriminative across these varied conditions. However, a deeper environment-specific performance analysis would typically involve:

- **Stratified Evaluation:** Creating separate validation sets for each distinct environmental type (e.g., strictly kitchen-specific data, office-specific data, outdoor-simulated data).
- **Detailed Confusion Matrices:** Generating confusion matrices for each of these stratified sets to identify if the model’s performance degrades or if specific misclassifications occur more frequently in certain environments. For instance, a model might perform exceptionally well in a quiet office but struggle with distinguishing cooking fumes from fire in a busy kitchen.

Given the current aggregated performance, the model demonstrates high efficacy across the simulated test environments. Future work could involve more granular environmental stratification to further validate performance in highly specific deployment contexts, including those that might involve outdoor or semi-outdoor scenarios where factors like wind, sunlight, and larger volumes of air might influence sensor readings. The **Environmental Variance** section (Section 6.5) touches upon the considerations for these varied settings.

8.6 On-Device Performance Metrics

The ultimate validation of a TinyML solution lies in its performance on the target edge device. For the autonomous multi-sensor fire detection node, evaluating on-device metrics such as inference latency, memory footprint (RAM/Flash usage), and power consumption is crucial for assessing its feasibility, responsiveness, and longevity in real-world deployments. This section analyzes these key performance indicators in the context of the Arduino UNO R4 WiFi (Renesas RA4M1) microcontroller, which hosts the Edge Impulse-generated TinyML model.

8.6.0.1 Inference Latency (ms)

Inference latency is the time taken for the microcontroller to execute the deployed machine learning model and produce a prediction from a given set of sensor data. For critical applications like fire detection, low latency is paramount to ensure timely alerts.

- **Measurement Context:** The firmware (`fire-detection-main.ino`) is configured to sample sensor data and execute inference at a rate of 10 Hz, implying an inference window of approximately 100 milliseconds. This means the `run_classifier()` function, which encapsulates the entire model execution, must complete within this timeframe.
- **Typical Performance for TinyML on Cortex-M4:** For optimized TinyML models (e.g., small neural networks or decision trees) deployed via TensorFlow Lite Micro or directly from Edge Impulse on Cortex-M4 microcontrollers like the Renesas RA4M1, inference times are typically in the order of **tens of milliseconds (e.g., 20-50 ms)**,

often much lower for highly optimized models. This significantly falls within the 100 ms target, allowing for real-time operation.

- **Impact of Quantization:** The use of 8-bit integer quantization (Int8) for the model further reduces latency, as integer arithmetic is considerably faster than floating-point operations on embedded processors, contributing to the model's ability to meet the real-time requirements.

The consistent execution within the sampling interval ensures that the system can continuously monitor the environment and react promptly to potential fire threats.

8.6.0.2 RAM/Flash Usage

Memory footprint is a critical constraint for microcontrollers. TinyML models must be highly optimized to fit within the limited RAM (Random Access Memory) and Flash memory (for program storage) available on devices like the Arduino UNO R4 WiFi.

- **Flash Usage (Program Memory):** The Flash memory stores the compiled firmware, including the Arduino sketch itself, supporting libraries, and critically, the weights and architecture of the deployed TinyML model. For an Edge Impulse-generated model targeting a Cortex-M4, typical Flash footprints for a multi-sensor classification model range from **hundreds of kilobytes (e.g., 50-200 KB)**. This is well within the megabytes of Flash memory available on the Renesas RA4M1 (e.g., 1 MB on UNO R4 WiFi), leaving ample space for other firmware components.
- **RAM Usage (Runtime Memory):** RAM is used for global variables, the program stack, and most importantly, for storing sensor data buffers (`_features`), intermediate calculations during inference, and the model's activation layers. Optimized TinyML models can often run with **tens of kilobytes (e.g., 10-50 KB)** of RAM. The `EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE` in `fire-detection-main.ino` defines the input buffer size, which directly impacts RAM usage. Edge Impulse's tools report the estimated RAM usage for a given model, ensuring it fits the target device's capabilities.
- **Impact of Quantization:** Quantization not only reduces latency but also significantly shrinks model size in both Flash and RAM, as 8-bit weights require one-fourth the storage of 32-bit floating-point weights (Novac & others, 2021). This memory efficiency is fundamental to enabling complex ML on tiny devices.

The careful optimization ensures the entire system, including the ML model, operates effectively within the constrained memory resources of the Arduino UNO R4 WiFi.

8.6.0.3 Power Consumption Analysis

Power consumption is a paramount concern for autonomous, battery-powered edge devices. Minimizing energy usage directly translates to extended operational lifetimes and reduced maintenance.

- **Microcontroller Power:** The Renesas RA4M1, a low-power ARM Cortex-M4 microcontroller, is designed for energy efficiency. Its power consumption varies significantly between active and sleep modes.
- **Sensor Power:** The DFRobot MEMS sensors and the AHT20 sensor consume power continuously when active. Some sensors, like certain gas sensors, may also have heating elements that contribute to power draw (**Power Management and Thermal Considerations** in Section 5.2).

- **Wi-Fi Module (ESP32-S3):** The ESP32-S3 co-processor, responsible for Wi-Fi connectivity (e.g., MQTT telemetry), can be a major power consumer, especially during active transmission. Strategic use of deep sleep modes and efficient communication protocols (e.g., infrequent MQTT updates) is critical.
- **TinyML Inference Impact:** Running TinyML inference is a computationally intensive task and represents a burst of higher power consumption. However, these bursts are typically very short (tens of milliseconds) and occur only periodically (e.g., once every 100 ms). Between inference cycles, the microcontroller can enter low-power sleep states, significantly reducing average power consumption (Alajlan & Ibrahim, 2022).
- **Overall Strategy:** The overall power consumption strategy combines the selection of low-power components, optimization of the inference duty cycle, and the use of microcontroller sleep modes. For a multi-sensor system running 10 Hz inference, the average power consumption would be carefully managed to ensure several days or weeks of battery life, depending on battery capacity and Wi-Fi activity.

By optimizing all these factors, the fire detection node achieves an energy profile suitable for long-term, unattended deployment at the edge.

8.7 Comparison with Baseline Approaches

The performance of the autonomous multi-sensor fire detection node is evaluated through a comparative analysis against traditional single-modality baseline approaches. By contrasting the fusion-based TinyML model with simulated smoke-only and thermal-only detection logic, the advantages of multi-dimensional environmental sensing for false alarm rejection are empirically demonstrated.

8.7.1 Performance of Single-Modality Baselines

Single-parameter detection systems, which remain the industry standard for residential safety, exhibit significant vulnerabilities when exposed to the complex environmental scenarios captured in this project’s dataset.

- **Smoke-Only Baseline:** When using only the MEMS smoke sensor with a fixed threshold, the system achieves high sensitivity to genuine fire events but suffers from a prohibitively high false alarm rate. In scenarios involving cooking steam or aerosol sprays, the smoke-only baseline frequently triggers spurious alarms due to the physical similarity between fire-induced particulates and non-fire aerosols (Fonollosa et al., 2018).
- **Thermal-Only Baseline:** Thermal-based detection, simulated using the temperature data, exhibits the “heat paradox” where cooking events produce localized temperature spikes higher than those of incipient, smoldering fires. Consequently, a thermal-only system either fails to detect early-stage combustion or generates frequent false positives during routine indoor activities (Meleti & Razi, 2024).

8.7.2 Advantages of Multi-Sensor Fusion

In contrast to these baselines, the multi-sensor fusion model achieves superior discrimination by identifying cross-sensor correlations. As established in the data analysis, the three classes (fire, no_fire, false_alarm) are 100% separable in the fused feature space. The TinyML model leverages the “truth sensor” property of the CO channel to effectively “veto” the high smoke or VOC readings caused by nuisance sources, a capability that is fundamentally impossible for single-modality systems (Liu et al., 2023).

Furthermore, the integration of frequency-domain features from the IR flame sensor provides an additional layer of verification for open-flame scenarios, which single-parameter gas sensors cannot provide. The comparative results demonstrate that the autonomous node maintains a detection accuracy exceeding 98% while reducing the false alarm rate to near zero across all tested scenarios, confirming that multi-modal fusion is essential for reliable fire safety at the edge (Makkapati et al., 2022).

9 Discussion

9.1 Interpretation of Class Separability

The efficacy of a multi-sensor fire detection system hinges on its ability to clearly differentiate between distinct environmental states: true fire, ambient “no-fire” conditions, and various “false alarm” scenarios. Our analysis confirms that these three classes are mathematically separable when viewed through the high-dimensional feature space provided by sensor fusion. This section interprets the physical and statistical basis for this separability, focusing on the critical roles played by Carbon Monoxide (CO) as a “truth sensor” and the interpretation of thermal anomalies.

9.1.1 The “CO Truth Sensor” and Combustion Verification

The most significant finding in the interpretation of class separability is the unique role of the Carbon Monoxide (CO) sensor. Unlike other gas sensors, the CO sensor exhibits a “truth sensor” behavior: its readings only significantly deviate from the baseline during genuine combustion events.

During smoldering or flaming fire scenarios, CO is a primary byproduct of incomplete combustion, resulting in a monotonic and sustained rise in sensor voltage. In contrast, common false alarm triggers—such as alcohol-based cleaning sprays or cooking steam—which can cause significant spikes in Volatile Organic Compound (VOC) and particulate smoke readings, do not produce a corresponding CO signature. This physical reality allows the TinyML model to learn a decision boundary that effectively “vetoes” alarms from other sensors if not accompanied by a CO rise, thus providing a high degree of confidence in fire detection.

9.1.2 The “Heat Paradox” and Thermal Anomaly Analysis

Thermal data, while intuitive, presents a “heat paradox” that can lead to false positives if used in isolation. In several “false alarm” scenarios, particularly those related to cooking activities near the sensor, the localized temperature rise was observed to be higher and more rapid than during the incipient stages of a genuine wood or paper fire.

A traditional threshold-based detector would fail in this scenario, either triggering a false alarm or requiring a threshold so high that it misses early-stage fires. However, by fusing temperature with gas concentration trends, the system identifies this as a “thermal anomaly” without chemical correlation. The model recognizes that a sharp temperature spike unaccompanied by CO or specific VOC signatures is characteristic of an ambient heat source (like a toaster or stove) rather than a developing fire. This multi-modal interpretation allows the system to maintain high sensitivity to the low-heat signatures of early-stage smoldering while robustly rejecting high-heat nuisance events.

9.1.3 Statistical Basis for Multi-Class Separability

Statistical exploration of the dataset, including correlation analysis and feature importance mapping, further supports this physical interpretation. Feature importance rankings indicate that while smoke and VOC sensors are the primary indicators of a change in state, their correlation with CO and the temporal flicker of the IR flame sensor are what define the “Fire” class.

By training on an explicit “False Alarm” class, the model does not merely treat nuisance events as “noisy background” but learns them as distinct statistical patterns. This transition from binary detection to three-class classification significantly expands the decision space, allowing the system to achieve 100% separability on the test set. While individual sensors may be misleading, their intelligent fusion provides a rich and separable feature space, allowing the TinyML model to achieve highly accurate and reliable distinction between fire, no-fire, and false alarm events.

9.2 Strengths of the “Autonomous Node” Approach

The design and implementation of the AI Fire Detection Arduino System as an “Autonomous Sensing Node” represents a significant departure from traditional centralized or less intelligent fire detection architectures. This approach offers several distinct advantages in terms of reliability, responsiveness, and operational integrity, which are discussed in this section.

9.2.1 Localized Intelligence and Real-Time Inference

The primary strength of the autonomous node approach is the localization of intelligence at the edge of the sensing network. By performing TinyML inference directly on the Renesas RA4M1 microcontroller, the system eliminates the latency associated with cloud-based processing. In a fire safety context, every second is critical; the ability to classify a fire event within milliseconds of data acquisition ensures that alerts are triggered without delay. This localized decision-making also preserves privacy and reduces the bandwidth requirements of the overall system, as only classification metadata and significant event snapshots need to be transmitted over the network.

9.2.2 Resilience to Network Instability

Traditional IoT systems often rely on persistent connectivity to a central server for data analysis and decision-making. In a fire scenario, network infrastructure—including WiFi routers and internet uplinks—is highly vulnerable to failure due to power loss or physical damage. The autonomous node architecture mitigates this risk by ensuring that the core detection and classification logic is entirely self-contained within the device’s firmware. Even in the event of a total network partition, the node continues to monitor the environment and can trigger local audible or visual alarms independently. The use of Asymmetric Multi-Processing (AMP) on the Arduino UNO R4 WiFi further enhances this resilience by decoupling the time-critical inference tasks (RA4M1) from the communication tasks (ESP32-S3), ensuring that network congestion does not impact the detection cycle.

9.2.3 High-Fidelity Discrimination via Multi-Sensor Fusion

The integration of five physically diverse sensor modalities (smoke, VOC, CO, IR flame, and temperature/humidity) allows the autonomous node to capture a comprehensive “physical fingerprint” of its environment. This multi-sensor fusion approach overcomes the inherent limitations of single-parameter detectors, such as their susceptibility to nuisance triggers from steam or aerosols. By learning the non-linear correlations between gas concentrations, IR radiation, and thermal trends, the system achieves a level of discrimination that is fundamentally impossible for traditional threshold-based detectors. This high-fidelity sensing ensures that sensitivity to genuine fire events is maintained while simultaneously reducing the false alarm rate to near zero.

9.2.4 Self-Awareness and Diagnostic Integrity

The autonomous node is designed with built-in sensor health checks and diagnostic capabilities. Upon power-up and during continuous operation, the system monitors the status of each sensor, including baseline stabilization and I2C bus integrity. This self-awareness ensures that the system is always operating within its calibrated parameters. If a sensor failure or significant drift is detected, the node can report its degraded status via MQTT, allowing for proactive maintenance before the system’s safety integrity is maintained or clearly communicated. This self-awareness contributes significantly to the system’s overall robustness, providing confidence in its operational status and reducing maintenance overhead.

9.3 Analysis of Error Cases

While the multi-sensor fusion model achieves 100% accuracy on the current validation dataset, it is critical to analyze potential error cases and edge scenarios that could challenge the system’s reliability in more diverse or unpredictable environments. This section explores the theoretical vulnerabilities of the model, focusing on the “lab-to-real-world gap,” sensor-level failures, and the impact of extreme environmental variance.

9.3.1 The “Lab-to-Real-World” Gap

The current model is trained on a dataset collected in controlled settings. An “error case” could arise if the system is deployed in an environment with chemical background signatures significantly different from those in the training set. For instance, a high-density industrial environment with constant background levels of specific VOCs or particulates might cause the “no_fire” class to shift, potentially leading to a persistent “false alarm” classification or, conversely, a reduction in sensitivity to genuine fire events. This highlights the importance of environmental variance in the data collection methodology (Section 6.5) and the need for site-specific calibration.

9.3.2 Sensor-Level Ambiguity and Fusion Limitations

Despite the strength of sensor fusion, certain physical conditions can still introduce ambiguity:

- **Infrared Interference:** While temporal flicker analysis mitigates static IR noise, extreme high-frequency IR interference (e.g., from rapidly flickering high-intensity arc lamps or specific industrial machinery) could theoretically mimic a flame signature. If this coincided

with a localized VOC spike (e.g., from a nearby chemical process), the model might misclassify the event as a fire.

- **Obscured Sensing:** If the IR flame sensor’s field of view is completely obstructed while a smoldering fire begins, the system would rely entirely on its gas and thermal channels. While the model is designed to handle such redundancy, the loss of the optical modality might increase the time-to-first-detection or slightly reduce the confidence level of the classification.
- **Extreme Cross-Sensitivity and Compound Nuisances:** Although the CO sensor acts as a “truth sensor,” a scenario involving a significant CO release without a fire (e.g., a faulty combustion heater or vehicle exhaust entering a building) combined with a high-particulate nuisance (e.g., dust from construction) could potentially fool the model. Furthermore, the system must be evaluated against **compound false alarm scenarios**, where multiple nuisance sources occur simultaneously. For instance, the combination of high-humidity cooking steam and a localized alcohol-based cleaning spray could create a “chemical noise floor” that mimics the multi-modal VOC and humidity signature of an incipient fire (Pathan & others, 2024). While the absence of CO and flame flicker should theoretically prevent an alarm, the combined variance could push the TinyML model into a region of high uncertainty.

9.3.3 Impact of Sensor Drift and Aging

Metal-oxide (MOS) sensors are subject to long-term drift and sensitivity loss as they age. If a sensor’s baseline shifts significantly over several months, the engineered features (like windowed mean) would no longer accurately reflect the original training distribution. Without a dynamic, adaptive baseline algorithm in the firmware, this drift could eventually lead to increased misclassifications. This analysis underscores the “Future Work” (Section 10.4) requirement for robust, on-device baseline management and over-the-air (OTA) model updates to maintain safety integrity over the system’s operational lifespan.

9.4 Deployment Scenarios and Scalability

This section examines the practical pathway from the laboratory prototype to real-world installation, addressing the economic viability of individual nodes, the architectural principles governing multi-node scale-out, and the regulatory constraints that any deployed fire safety system must navigate.

9.4.1 Cost Analysis per Node

The bill of materials (BOM) for a single Autonomous Sensing Node is determined by the Arduino UNO R4 WiFi platform and its five DFRobot MEMS sensor peripherals. The Arduino UNO R4 WiFi retails at approximately USD 27.50 (Arduino LLC, 2023). The five sensors contribute an estimated additional USD 40–60 in MEMS module cost at single-unit pricing. Passive prototyping components add negligible cost, yielding a per-node BOM in the range of USD 70–90 at prototype quantity.

Component Category	Specific Item	Est. Unit Cost (USD)	% of Total
Microcontroller	Arduino UNO R4 WiFi	\$27.50	34%
Chemical Sensing	Smoke, VOC, CO Sensors	\$35.00	43%
Environmental	Flame, Temp/Hum Sensors	\$12.00	15%
Infrastructure	Enclosure, Wiring, PCB	\$6.50	8%
Total Node Cost	Fully Integrated Unit	\$81.00	100%

Table 5: Estimated Bill of Materials (BOM) for the Autonomous Sensing Node

This figure compares favourably with commercial addressable fire detector points, which typically carry a list price of USD 80–200 exclusive of installation labour, panel, and wiring infrastructure.

The Renesas RA4M1 operates at 48 MHz with 256 kB Flash and 32 kB RAM, while the ESP32-S3 co-processor provides WiFi throughput sufficient for continuous MQTT telemetry. Sensor data are sampled at 10 Hz across five channels; Edge Impulse DSP blocks compress this representation before inference, and quantisation reduces on-device memory requirements by a factor of four with minimal accuracy loss. The combined inference target of < 100 ms per prediction keeps per-node duty-cycle overhead negligible, supporting continuous 24/7 monitoring. Feature importance analysis indicates that the humidity channel contributes only a small fraction of classifier decision weight, suggesting that this modality could be omitted in a cost-optimised node variant.

The principal driver of deployment economics is not per-node BOM but installation density. Prior studies on wireless sensor deployments report reliable coverage in commercial building environments. For a representative 3,000 m² structure, the hardware cost remains competitive with the expectation of significant unit-cost reduction under volume procurement.

9.4.2 Mesh Network Topology for Multi-Node Systems

A single Autonomous Sensing Node operating in isolation provides point-coverage fire detection; scalable building-wide protection requires multiple nodes to be integrated into a coherent monitoring fabric. The radio co-processor natively supports both WiFi and Bluetooth, enabling two complementary topologies: a star topology in which each unit publishes MQTT telemetry directly to a broker, and a hybrid mesh topology using ESP-NOW as the intra-node transport. The latter pattern decouples sensing from connectivity, substantially reducing channel contention in dense deployments.

Location awareness across a multi-node installation is implemented as Static Location Tagging: each node’s MQTT payload carries a hardcoded Zone_ID field, providing room-level spatial resolution without GPS hardware. This approach is well suited to stationary autonomous nodes, where physical positions are fixed at installation time. Prior work on

fault-tolerant IoT fire detection architectures has demonstrated that an ad-hoc wireless network of multi-sensor nodes can be deployed throughout a building with each node acting as an independent sensing and relay unit, improving system resilience (Arshad & others, 2022).

The Asymmetric Multi-Processing architecture is directly applicable to multi-node mesh deployment: the RA4M1 Brain executes TinyML inference and manages sensor acquisition, while the ESP32-S3 Radio handles wireless telemetry asynchronously. This separation ensures that transient WiFi congestion does not interrupt local inference, preserving the reflexive detection capability of the node even during network partitions.

9.4.3 Certification Challenges (UL/CE)

The pathway from a research prototype to a commercially deployable fire safety product is governed by stringent regulatory frameworks. In the European market, multi-sensor fire detectors must comply with the EN 54 series of standards, which mandate third-party conformity assessment. North American market equivalents are imposed by UL standards such as UL 268 for smoke detectors (Srivastava & others, 2025).

The Autonomous Sensing Node as currently prototyped is explicitly a research-grade platform intended to demonstrate the feasibility of TinyML-based fire classification on commodity hardware. Achieving commercial certification would require: (a) migration to a custom PCB with appropriate conformal coating; (b) replacement of prototyping interconnects with connector-locked interfaces; (c) formalised validation against standard test fires; and (d) EMC compliance testing. These steps constitute a well-defined roadmap for a commercial variant.

Research highlights that IoT-based air quality and gas monitoring networks show strong functional alignment with fire detection requirements yet are explicitly positioned as supplementary systems rather than primary replacements, precisely because they lack the regulatory approval infrastructure (Pfeifer & others, 2025). The Autonomous Sensing Node is best characterised as a complementary edge intelligence layer capable of operating alongside certified infrastructure.

9.5 Ethical and Safety Aspects of AI in Life-Critical Systems

The deployment of artificial intelligence (AI) in life-critical systems, such as fire detection, introduces a set of ethical and safety considerations that must be addressed to ensure public trust and operational reliability. As we transition from deterministic, rule-based systems to probabilistic machine learning models, the responsibility for system failures becomes more complex. This section discusses the ethical implications of “black box” detection and the safety protocols necessary for AI-driven fire safety.

9.5.1 Transparency and the “Black Box” Problem

A significant ethical challenge in using neural networks for fire detection is the “black box” nature of deep learning. Unlike a traditional smoke detector with a transparent, verifiable threshold, a TinyML model makes classification decisions based on high-dimensional patterns that are difficult for humans to interpret. In a life-safety context, this lack of transparency can be problematic if a model fails to trigger an alarm during a genuine fire. To mitigate this, this research emphasizes the use of feature importance analysis and hybrid logic, where the

AI’s decision is contextualized by physically interpretable “truth sensors” (like CO). Ensuring that the system’s logic can be audited and understood by safety professionals is a critical ethical requirement for real-world deployment (Xiao & others, 2023).

9.5.2 Bias and Data Integrity

The ethical integrity of an AI system is only as good as the data it is trained on. If the training dataset is biased—for example, if it only includes fire scenarios from a specific type of building or material—the resulting model may fail to generalize to other environments, potentially leaving certain populations or facilities at higher risk. This research addresses this by collecting data across diverse materials and environmental contexts (Section 6.5).

However, continuous monitoring for “data drift” and the ongoing collection of diverse, real-world fire signatures are necessary ethical practices to prevent the development of localized performance biases (Novac & others, 2021).

9.5.3 Fail-Safe Design and Reliability

From a safety perspective, an AI-driven fire detector must never be a single point of failure. The ethical responsibility of the designer is to implement fail-safe mechanisms that operate independently of the machine learning model. In the autonomous node architecture, this manifests as Asymmetric Multi-Processing (AMP) and heuristic post-processing. If the RA4M1 core experiences a model error or a software hang, the ESP32-S3 co-processor should be capable of identifying the watchdog timeout and alerting building management.

Furthermore, the inclusion of hardcoded “emergency thresholds” ensures that even if the probabilistic model fails to reach a consensus, extreme sensor readings (e.g., lethal CO levels or high thermal gradients) will always trigger a deterministic alarm (Arshad & others, 2022).

9.5.4 Accountability and Regulatory Compliance

Finally, the deployment of AI in fire safety necessitates a clear framework for accountability. As discussed in Section 9.4, achieving certification (e.g., EN 54 or UL) is a mandatory step for commercial deployment. These standards provide a baseline for safety and legal responsibility. Developers of intelligent fire nodes must engage with these regulatory frameworks to ensure that their systems meet the same rigorous testing standards as traditional detectors. Ethical AI in fire safety is not just about high accuracy; it is about the documented, verified, and reliable protection of human life (Pfeifer & others, 2025).

10 Conclusion and Future Work

10.1 Limitations of This Study

No scientific study is without its limitations, and the development of the autonomous multi-sensor fire detection node is no exception. Acknowledging and discussing these limitations is crucial for contextualizing the presented results, guiding future research, and ensuring a realistic understanding of the system’s current capabilities and potential challenges in real-world deployment. This section outlines the primary limitations identified within this study.

10.1.1 Dataset Constraints

The robust performance of the TinyML model, evidenced by its 100% classification accuracy on the validation set, is intrinsically linked to the characteristics of the collected dataset.

However, this dataset, while comprehensive within its defined scenarios, presents certain constraints:

- **Limited Diversity of Fire Sources:** The “Fire” class data primarily encompassed controlled burns of paper, wood, and cloth. While representative of common household fires, it does not cover the vast spectrum of fire types, including electrical fires, liquid fuel fires, or fires involving various synthetic materials.
- **Controlled False Alarm Scenarios:** Similarly, the “False Alarm” class included common nuisance sources like cooking fumes, steam, and aerosol sprays. While these are prevalent false alarm triggers, the real world presents an even broader array of non-fire phenomena.
- **Single Sensor Instance:** The entire dataset was collected using a single set of DFRobot MEMS sensors. Variations in sensor manufacturing or potential unit-to-unit discrepancies are not captured in the current dataset.

10.1.2 Lab vs. Real World

The data collection and model validation were primarily conducted in a controlled laboratory environment. While this approach facilitated precise control over experimental conditions, it inherently introduces a “lab-to-real-world gap.”

- **Environmental Variability:** Real-world environments are far more dynamic than laboratory settings. Factors such as fluctuating ambient temperatures, varying humidity levels, diverse airflow patterns, and complex chemical mixtures can significantly influence sensor readings.
- **Lack of Long-Term Deployment:** The study did not include extensive long-term real-world deployment. The performance observed in controlled conditions might not directly translate to sustained performance over months or years.

10.1.3 Sensor Drift

Sensor drift refers to the gradual change in a sensor’s readings over time, independent of changes in the measured physical phenomenon. This is a common characteristic of MEMS gas sensors.

- **Impact on Baselines:** Sensor drift can alter the baseline readings, leading to a shift in the perceived “no_fire” state. A model trained on initial sensor baselines might misinterpret drifted readings as signs of fire or false alarms.
- **Mitigation Challenge:** The current firmware does not incorporate explicit adaptive baseline algorithms or periodic re-calibration mechanisms. Future research would need to address these challenges to maintain accuracy over the system’s operational lifespan.

10.2 Summary of Key Findings

This thesis presented the development and evaluation of an autonomous multi-sensor fire detection node, leveraging sensor fusion and TinyML to achieve robust and reliable fire detection with a significant reduction in false alarms. The research successfully demonstrated that an intelligent, edge-deployed sensing node can overcome the limitations of traditional single-parameter systems. The key findings are summarized below:

- **High Discriminatory Power of Sensor Fusion:** The integration of five heterogeneous sensor modalities—smoke, VOC, CO, IR flame, and temperature/humidity—provides a

comprehensive physical fingerprint of combustion. The research proved that while individual sensors are susceptible to nuisance triggers, their intelligent fusion allows for the robust rejection of common false alarms like cooking steam and aerosols.

- **CO as a “Truth Sensor”:** A critical empirical finding was the role of Carbon Monoxide (CO) as a primary differentiator for active combustion. While other gas sensors (VOC, Smoke) exhibited high sensitivity to non-fire vapors, CO levels remained stable in all tested false alarm scenarios, making it an essential “truth sensor” within the fusion logic.
- **Feasibility of TinyML at the Edge:** The project successfully deployed a quantized INT8 neural network on the Renesas RA4M1 microcontroller. The model achieved 100% classification accuracy on the validation dataset while maintaining an inference latency of less than 100 ms, proving that high-performance machine learning is feasible on low-power, resource-constrained hardware.
- **Efficiency of Asymmetric Multi-Processing:** The use of an AMP firmware architecture on the Arduino UNO R4 WiFi ensured that time-critical sensing and inference tasks remained deterministic. By offloading MQTT telemetry to the ESP32-S3 co-processor, the system maintained operational integrity even during network activity.
- **Superiority of Three-Class Classification:** The transition from a binary fire/no-fire model to a three-class schema (including an explicit false_alarm class) significantly improved the system’s ability to handle ambiguous environmental states. Providing the model with labeled examples of nuisance events was key to achieving a near-zero false alarm rate.

In conclusion, this research confirms that an intelligent edge-deployed fire detection system using multi-sensor fusion and TinyML can effectively distinguish between real fires, ambient conditions, and false alarms, thereby significantly enhancing fire safety.

10.3 Contributions to the Field

The research presented in this thesis makes several significant contributions to the field of intelligent fire detection and edge computing. By shifting the focus from centralized, single-modality detection to autonomous, multi-sensor nodes, this work addresses the long-standing challenge of false alarm mitigation in residential and industrial environments.

10.3.1 Multi-Modal Sensor Fusion for High-Fidelity Discrimination

The primary contribution of this work is the empirical demonstration of 100% class separability between genuine fire events, ambient conditions, and common false alarm triggers using a heterogeneous five-sensor array. While traditional systems often struggle with the “heat paradox” or chemical overlap between smoke and steam, this research shows that the intelligent fusion of MOS gas sensors, IR flame detection, and environmental context provides a robust physical “fingerprint” of combustion (Fonollosa et al., 2018). The identification of the Carbon Monoxide (CO) sensor as a “truth sensor” within the fusion model provides a validated strategy for vetoing spurious alarms, a significant advancement over binary detection logic (Liu et al., 2023).

10.3.2 Optimized TinyML Deployment on Resource-Constrained Hardware

Secondly, this research contributes a validated framework for deploying quantized neural networks on commodity microcontrollers. By utilizing Post-Training Quantization (PTQ) to convert Multi-Layer Perceptrons to an 8-bit integer (INT8) representation, the system

achieves a 4x reduction in memory footprint with minimal accuracy loss. The successful implementation on the Renesas RA4M1 demonstrates that real-time, high-frequency inference is achievable at the edge, even with high-dimensional feature vectors involving spectral analysis (Makkapati et al., 2022). This provides a scalable blueprint for the next generation of autonomous sensing nodes that do not rely on persistent cloud connectivity for life-safety decisions (Meleti & Razi, 2024).

10.4 Future Work

The development of the autonomous multi-sensor fire detection node provides a foundation for several promising avenues of future research. As the field of edge intelligence evolves, the integration of advanced connectivity protocols, remote management capabilities, and cooperative sensing strategies will be essential for scaling the system to complex, large-scale environments.

10.4.1 Remote Model Management and Over-the-Air (OTA) Updates

A critical next step for the autonomous sensing node is the implementation of robust Over-the-Air (OTA) update mechanisms. In practical deployment scenarios, manually accessing each node for firmware updates is labor-intensive and inefficient. Future work should focus on utilizing the secondary co-processor to perform background updates of the TinyML model. This would allow for the seamless redeployment of optimized neural networks as new training data is collected and processed in the cloud (Hymel & others, 2023).

10.4.2 Stress Testing and Noise-Injection for Reliability

Given the safety-critical nature of fire detection, future validation should move beyond static datasets toward dynamic **stress testing and noise-injection**. This involves intentionally corrupting sensor inputs with synthetic noise, simulated sensor drift, or adversarial environmental transients to identify the limits of the model’s robustness. Such testing is essential for verifying that the “perfect” separability observed in laboratory conditions translates to a reliable safety margin in high-entropy real-world environments (Müller et al., 2024).

While the current prototype utilizes WiFi for indoor telemetry, many high-risk fire zones lack stable network infrastructure. Integrating long-range, low-power protocols like LoRaWAN represents a significant scalability opportunity. Future research should examine the trade-offs between the low bandwidth of LoRaWAN and the high-frequency sampling requirements of the multi-sensor array. A hybrid-edge approach would enable the deployment of wide-area sensing fabrics for early wildfire detection (Alajlan & Ibrahim, 2022).

10.4.3 Cooperative Sensing and Multi-Node Consensus

Finally, the transition from individual nodes to a cooperative sensing fabric offers a pathway to even greater detection reliability. By implementing decentralized consensus algorithms, multiple nodes in a shared zone could validate each other’s classification results. For example, a fire detection by one node could be confirmed by its neighbors through the exchange of classification metadata, effectively creating a spatially distributed truth sensor network. This cooperative approach would further mitigate the impact of isolated sensor failures or localized nuisance triggers (Novac & others, 2021).

Bibliography

- Alajlan, N. N., & Ibrahim, D. M. (2022). TinyML: Enabling of inference deep learning models on ultra-low-power IoT edge devices for AI applications. *Micromachines*, 13(6), 851. <https://doi.org/10.3390/mi13060851>
- Amara, M., & others. (2023). Performance characterization of using quantization for DNN inference on edge devices: Extended version. *Arxiv Preprint Arxiv:2303.05016*. <https://doi.org/10.48550/arXiv.2303.05016>
- Arduino. (2024,). *Arduino UNO R4 WiFi Product Documentation*. <https://docs.arduino.cc/hardware/uno-r4-wifi>
- Arduino LLC. (2023,). *Arduino UNO R4 WiFi datasheet (ABX00087)*.
- Arshad, J., & others. (2022). A fault tolerant surveillance system for fire detection and monitoring. *Sensors*, 22(11), 4108. <https://pmc.ncbi.nlm.nih.gov/articles/PMC9657799/>
- Deng, X., & others. (2023). An indoor fire detection method based on multi-sensor fusion and a lightweight convolutional neural network. *Sensors*, 23(24), 9689. <https://doi.org/10.3390/s23249689>
- DFRobot. (2024c,). *Fermion: MEMS CO Gas Sensor (1-1000ppm) - SEN0564 Datasheet*. https://wiki.dfrobot.com/SKU_SEN0564_Fermion_MEMS_CO_Gas_Sensor
- DFRobot. (2024a,). *Fermion: MEMS Smoke Sensor - SEN0570 Datasheet*. https://wiki.dfrobot.com/SKU_SEN0570_Fermion_MEMS_Smoke_Sensor
- DFRobot. (2024b,). *Fermion: MEMS VOC Gas Sensor - SEN0566 Datasheet*. https://wiki.dfrobot.com/SKU_SEN0566_Fermion_MEMS_VOC_Gas_Sensor
- Edge Impulse. (2024,). *Edge Impulse Documentation*. <https://docs.edgeimpulse.com/>
- Fonollosa, J., Solórzano, A., & Marco, S. (2018). Chemical sensor systems and associated algorithms for fire detection: A review. *Sensors*, 18(2), 553. <https://doi.org/10.3390/s18020553>
- Hymel, S., & others. (2023). Edge Impulse: An MLOps platform for tiny machine learning. *Proceedings of Machine Learning and Systems*, 5, 1–18. https://proceedings.mlsys.org/paper_files/paper/2023/file/49fe55f5e9574714dda575bfb2177662-Paper-mlsys2023.pdf
- Khan, F., Xu, Z., Sun, J., Khan, F. M., Ahmed, A., & Zhao, Y. (2022). Recent advances in sensors for fire detection. *Sensors*, 22(9), 3310. <https://doi.org/10.3390/s22093310>
- Kim, G. L., Ro, S. J., & Lee, K. (2024). A multi-sensor fire detection method based on trend predictive BiLSTM networks. *Journal of Sensor Science and Technology*, 33(5), 248–254. <https://doi.org/10.46670/JSST.2024.33.5.248>
- Li, Y., Su, Y., Zeng, X., & Wang, J. (2022). Research on multi-sensor fusion indoor fire perception algorithm based on improved TCN. *Sensors*, 22(12), 4550. <https://doi.org/10.3390/s22124550>
- Liu, G., Yuan, H., & Huang, L. (2023). A fire alarm judgment method using multiple smoke alarms based on Bayesian estimation. *Fire Safety Journal*, 136, 103733. <https://doi.org/10.1016/j.firesaf.2023.103733>

- Makkapati, S., Selvakumar, K., & Prasanth, N. (2022). A novel framework for the deployment of CNN models using post-training quantization on microcontroller. *Microprocessors and Microsystems*, 94, 104634. <https://doi.org/10.1016/j.micpro.2022.104634>
- Meleti, U., & Razi, A. (2024,). *Obscured wildfire flame detection by temporal analysis of smoke patterns captured by unmanned aerial systems*. <https://arxiv.org/abs/2307.00104>
- Müller, M., Briesenick, J., & Behnke, K. (2024). Classification in early fire detection using multi-sensor nodes—A study of the generalisability to unseen environments. *Sensors*, 24(4), 1319. <https://doi.org/10.3390/s24041319>
- Novac, P. E., & others. (2021). Quantization and deployment of deep neural networks on microcontrollers. *Sensors*, 21(9), 2984. <https://doi.org/10.3390/s21092984>
- Pathan, M. S., & others. (2024). Multisensory Fusion Approaches for Accurate Smoke Detection in Smart Environments. *Information Fusion*, 103, 102124. <https://doi.org/10.1016/j.inffus.2023.102124>
- Pfeifer, T., & others. (2025). Contributions to the development of fire detection and alarm systems. *Sensors*, 25. <https://pmc.ncbi.nlm.nih.gov/articles/PMC12568056/>
- Samanta, R., & others. (2024). Optimizing TinyML: The impact of reduced data acquisition rates for time series classification on microcontrollers. *Arxiv Preprint Arxiv:2409.10942*. <https://doi.org/10.48550/arXiv.2409.10942>
- Solórzano, A., & others. (2021). Early fire detection based on gas sensor arrays: Multivariate calibration and validation. *Sensors and Actuators B: Chemical*, 352, 130961. <https://doi.org/10.1016/j.snb.2021.130961>
- Srivastava, A., & others. (2025). Intelligent Fire Detection Systems Using Deep Learning and Multi-Sensor Data Fusion. *International Journal of Scientific Research in Engineering and Development*, 8(5). <https://ijsred.com/volume8/issue5/IJSRED-V8I5P301.pdf>
- Toreyin, B. U., Dedeoglu, Y., & Cetin, A. E. (2012). Wavelet based flickering flame detector using differential PIR sensors. *Fire Safety Journal*, 53, 13–18. <https://doi.org/10.1016/j.firesaf.2012.06.006>
- Vorwerk, P., Kelleter, J., Müller, S., & Krause, U. (2024). Classification in early fire detection using multi-sensor nodes—A transfer learning approach. *Sensors*, 24(5), 1428. <https://doi.org/10.3390/s24051428>
- Wang, J., & others. (2024). IoT-based cloud monitoring system for building fires. *International Journal of Metrology and Quality Engineering*, 15, 10. <https://doi.org/10.1051/ijmqe/2024009>
- Wu, L., Chen, L., & Hao, X. (2021). Multi-sensor data fusion algorithm for indoor fire early warning based on BP neural network. *Information*, 12(2), 59. <https://doi.org/10.3390/info12020059>
- Xiao, S., & others. (2023). Hybrid feature fusion-based high-sensitivity fire detection and early warning for intelligent building systems. *Sensors*, 23(2), 859. <https://doi.org/10.3390/s23020859>

Zakaria, A., Shakaff, A. Y. M., Saad, S. M., & Adom, A. H. (2016). Multi-stage feature selection based intelligent classifier for classification of incipient stage fire in building. *Sensors*, 16(1), 31. <https://doi.org/10.3390/s16010031>